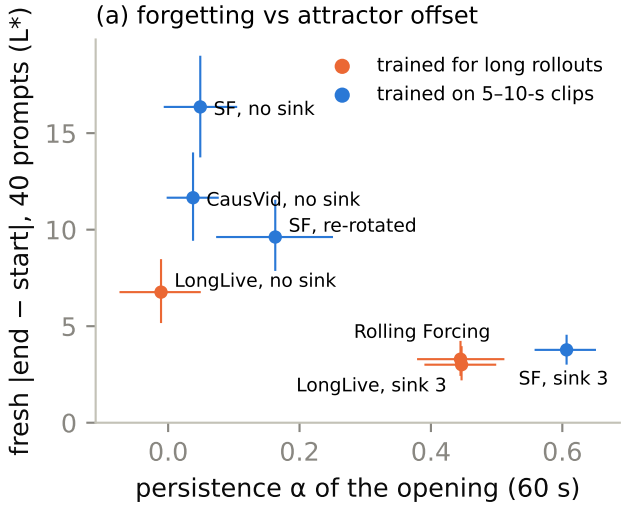
(a) forgetting vs attractor offset
trained for long rollouts
trained on 5–10-s clips
SF, no sink
CausVid, no sink
SF, re-rotated
LongLive, no sink
Rolling Forcing
LongLive, sink 3
SF, sink 3
fresh |end − start|, 40 prompts (L*)
persistence α of the opening (60 s)
0
5
10
15
0.0
0.2
0.4
0.6

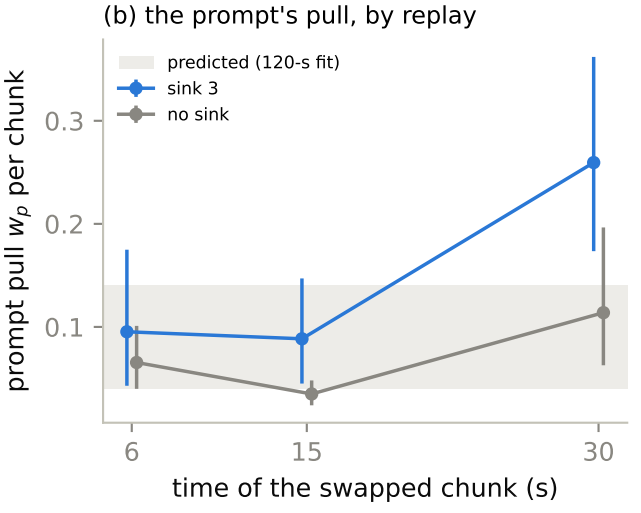
(b) the prompt's pull, by replay
predicted (120-s fit)
sink 3
no sink
prompt pull w_p per chunk
time of the swapped chunk (s)
0.1
0.2
0.3
6
15
30

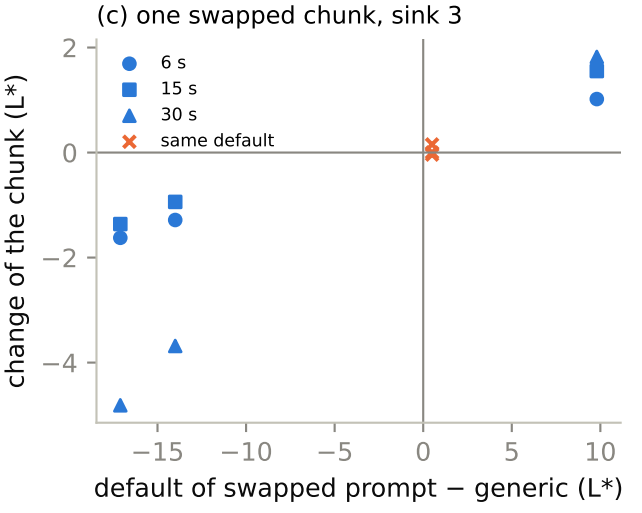
(c) one swapped chunk, sink 3
6 s
15 s
30 s
same default
change of the chunk (L*)
default of swapped prompt − generic (L*)
2
0
−2
−4
−15
−10
−5
0
5
10