Drift to Default: How Autoregressive Video Rollouts Forget Their Opening

Anonymous authors

Abstract

Autoregressive video diffusion models generate long videos chunk by chunk and drift over tens of seconds. We measure drift by intervention: we shift the brightness or colour of a held opening frame within a scene, or start from openings of different scene classes, and measure how much of the opening survives. In 36 preregistered studies on nine open models (1.3B to 14B parameters, three model families), rollouts *relax*: global tone and scene content move from the opening toward a default set by the conditioning, keeping a share α that depends on the model, the prompt and the horizon. Under a generic city prompt, openings of six scene classes lose 39–93% of their class evidence within 30 s, depending on the configuration. How much is kept depends on whether the opening stays in the model's context. Once it leaves, it is forgotten, through a sliding key–value window at 1.3B and at 14B parameters, and after a delay of about 20 s when overlapping frames are re-denoised. An opening that stays in the context keeps a share whose size depends on the model: 0.16–0.5 of its tone at 60 s for multi-frame sinks, anchors and clean prefixes (a re-rotated sink still keeps 0.10 at 240 s, and an anchor a steady 0.3), while a single anchored latent frame keeps almost no tone (0.03–0.08) but, in the 14B model, a third of the scene. Attention sinks left at their absolute rotary positions keep more, much of it in bursts, brief copies of the opening at the query–sink distances where an attention probe finds the sink most attended; they move when the sink is moved and recur at the times the probe predicts, out of sample to 180 s, and vanish in two models when one of the sink's 22 rotary frequencies is held. Drift separates into forgetting the opening and an offset of the default from where the same prompt starts; over 40 prompts the offset depends on inference settings as much as on training: it grows with the key–value window up to 21 frames in Self Forcing (5 to 16 L*) and LongLive, though not in CausVid, and not significantly beyond. A content-aware correction and a deeper sink reduce drift without training. We report all 125 preregistered primary tests, 34 of which failed.

1 Introduction

Autoregressive video diffusion models generate a video chunk by chunk, each chunk conditioned on the model's own earlier output [3, 1, 2]. This makes streaming and open-ended generation possible and is the basis of interactive world models [38, 39, 40]. It also makes them drift: over tens of seconds the picture darkens or saturates, loses texture, and its content wanders. A growing family of methods targets this failure with rolling windows, attention sinks, re-aligned positions, longer self-rollouts and test-time correction [7, 8, 9, 22, 12, 13, 21]. Most report a few samples per prompt and an aggregate quality score. What drift *is*, in quantitative terms, is rarely measured.

We measure it by intervention. We hold an opening frame, shift its brightness or colour *within* a scene, let the model continue, and ask how much of the shift survives; for scene content we start from openings of different classes. In every model we tested the rollout *relaxes*: it moves from its opening toward a default state set by the conditioning and keeps a share α of where it started that depends on the model, the prompt and the horizon,

$$x(T) \approx \alpha(T)\, x_{\text{open}} + \big(1 - \alpha(T)\big)\, \mu. \tag{1}$$

This is a linear-response summary, not a law: α and μ come from one regression. What it organises are three findings, and tests of its assumptions (Section 4).

What survives depends on whether the opening stays in the context (Fig. 3). Once the opening leaves a sliding key–value window it is forgotten, in 1.3B models and in a 14B model alike; re-denoising overlapping frames passes it forward and delays forgetting by about 20 s. An opening that stays in the context (a re-rotated sink, an anchor, a clean copy in every section, full history) keeps a share whose size depends

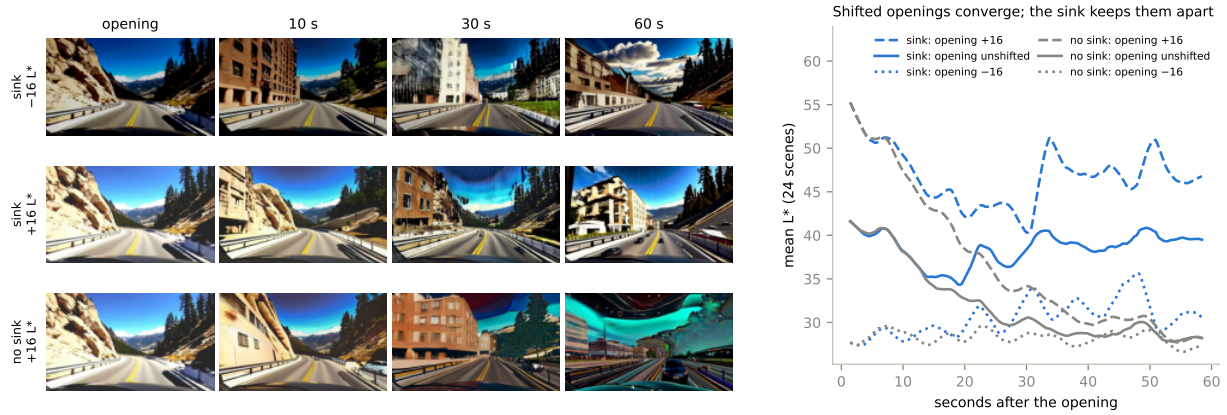


Figure 1: Relaxation. Left: a generated mountain-road opening, darkened or brightened by 16 L* before encoding, continued for 60 s under a generic city prompt by Self Forcing DMD with and without an attention sink. The scene turns into a city street in all three runs; without a sink the colours also blow out. Right: mean brightness over the 24 scenes of study 7 for the three versions of each opening (3-s moving average). Without a sink the versions converge to one default; with a sink they stay apart by a share that fluctuates around half.

on the model and, in the two we ran to 240 s, is still there then; tone and scene can be kept separately: a one-frame anchor in the 14B model keeps the scene but not its tone, LongLive's sink the reverse. A sink left at its absolute rotary positions keeps more, much of it in bursts in which the opening returns when the sink's distance to the current chunk passes a peak of the attention it draws; moving the sink moves the bursts, the probe predicts them out of sample to 180 s, and holding one rotary frequency of the sink removes them in two models.

For a given checkpoint, the default is set by the conditioning. Lighting words move the default brightness by up to 27 L*, and scene words decide whether a snowy, desert or forest opening stays or becomes a city street within 30 s. Which parts of the conditioning carry tone and which carry the scene differs between models.

Drift has two separable parts: forgetting the opening, $(1-\alpha)(x_{\text{open}}-\mu)$, and an offset of the default itself from where the same prompt starts, 3–16 L* over 40 prompts; it grows with the key–value window up to 21 frames in Self Forcing and LongLive but not CausVid, and with LongLive's window Self Forcing leaves no significant gap to LongLive (the preregistered test of a remaining gap failed). Training-free levers act on these terms: an early-step tone correction, a deeper sink and a sink held where the probe puts a moderate share, each with its own cost; lighting words matched to a model's measured defaults helped on one model only.

Contributions.

1. A measurement protocol for long-horizon drift: held and perturbed openings, persistence α and default μ for tone and scene content, paired seeds pinned to one GPU architecture, and preregistered analyses.
2. An empirical account of relaxation across nine open models (six of the Wan2.1-1.3B architecture, the Wan2.1-14B student Krea Realtime, the HunyuanVideo-based FramePack and MAGI-1), for tone and scene content, at horizons up to 240 s: what sets μ, what sets the long-run share, and a decomposition of drift into forgetting and an attractor offset, whose size we compare across models and key–value windows (preregistered).
3. A positional account of attention-sink bursts, predicted by an attention probe and tested by moving and holding the sink in two models.
4. Training-free levers, with their costs and failures: three that raise α (one placing the sink where the probe says), and a choice of prompt that moves μ on one model only.

We preregistered 36 studies with hashed analysis code and report all 125 primary tests, including the 34 that failed (Table 3). No model is trained; all code and per-video statistics are public.

2 Related work

Autoregressive and long video diffusion. Diffusion Forcing [3] and History-Guided Video Diffusion [33] enable causal rollout with per-frame noise and guidance on the history. CausVid [1] distils a bidirectional teacher [4] into a few-step causal student with DMD [5]; Self Forcing [2] and Self-Forcing++ [22] train on the student's own (minute-long) rollouts. LongLive [7], Rolling Forcing [8] and Deep Forcing [9] add attention sinks, rolling joint denoising and deeper, re-aligned sinks; LoL [21] traces "sink collapse" to positional periodicity [51, 52]. Segment-wise models condition on overlapping frames [10, 26, 31], a temporal pyramid [30] or packed history [23]; training-free schemes lengthen bidirectional models [25, 27, 28, 29]. These works propose mechanisms against drift; we measure what drift does and which mechanism changes which part of it.

Error accumulation, sinks and correction. Sampling from one's own output causes exposure bias in sequence models [34, 35] and diffusion models [36, 37]; error analyses of autoregressive video diffusion [14] bound how errors compound, and world models [38, 39, 40, 50] add noise to context frames for robustness. For global statistics our results suggest a complementary view: the rollout forgets its context and relaxes to the conditioning's default. Global brightness biases of diffusion models arise at the earliest denoising steps [41, 42]. Streaming language models keep the first tokens as an attention sink [11, 43], related to registers in vision transformers [44]; the closest analogue to our finding is instruction drift in dialogs [45], where attention to the system prompt decays and behaviour reverts to a default. Pathwise test-time correction [12] and FreqForcing [13] anchor rollouts to early frames or their low frequencies; our correction is a minimal member of this family. VBench [15] scores imaging quality with MUSIQ [16]; we measure drift against each prompt's fresh-start distribution [17] and by intervention.

3 Measuring drift by intervention

Models. We study nine open models (Appendix Table 2). Self Forcing (five checkpoints), CausVid, LongLive and Rolling Forcing are causal few-step students of Wan2.1-T2V-1.3B [4] that read their past through a key–value cache of 21 latent frames (LongLive: 12), optionally with an attention sink (the first chunk kept in the cache); Krea Realtime 14B [32] is a student of the 14B model. SkyReels-V2-DF [10] (same architecture as the 1.3B model) extends videos in overlapping segments, Matrix-Game-2.0 [50] is a Wan-architecture world model driven by actions, FramePack [23] (13B, from HunyuanVideo [24]) conditions every section on the opening, and MAGI-1 [31] (4.5B) has its own VAE and transformer; Section 5 describes how each holds the opening. All Wan-1.3B cache configurations run in one harness that reproduces the official Self Forcing loop bit for bit when no intervention is applied.

Held and perturbed openings. An opening image is held still for the first chunk (9 frames, 0.56 s at 16 fps; SkyReels: 17 frames as a prefix video; FramePack and Krea: the input image; MAGI-1: a clean image prefix) and continued under a prompt. Openings are NAVSIM [19] front-camera frames or frames generated by the models themselves. To separate what the model keeps from what the scene is, we shift the *same* opening in CIELAB before encoding: L^* by ± 8 and ± 16, b^* by ± 12. For scene content, we instead generate openings of six scene classes and vary the prompt (Section 4.1).

Persistence and default. For a statistic x (a per-frame CIELAB mean), the persistence at time T is the within-scene least-squares slope of the end value $\bar{x}(T)$, averaged over the window $[T-w, T)$ after the opening, on the opening value x_{open} (mean of the held frames),

$$\alpha(T) = \frac{\partial\, \bar{x}(T)}{\partial\, x_{\text{open}}}\Big|_{\text{scene}}, \qquad \mu(T) = \frac{b_0}{1-\alpha(T)}, \tag{2}$$

where b_0 is the fitted intercept, so μ is the level at which the end equals the opening (the zero-drift level). The window is 10 s for endpoints at 30 s or later (2–5 s in the earlier studies with shorter endpoints, 60 s in studies 29 and 31, 5 s for content) and 1 s for curves; Appendix Table 4 compares most models over one window, 15–20 s. Averaging slopes over scenes gives a scene-bootstrap interval. The fresh default is where prompt-only rollouts end.

Content persistence. For scene content we score each frame with CLIP ViT-L/14 [48] as the log-odds of the opening's class against "a city street", $z = s\,(\cos(f, t_{\text{class}}) - \cos(f, t_{\text{city}}))$, and define $\kappa(T) = (z(T) - \mu_z)/(z_{\text{open}} - \mu_z)$ with μ_z the class score of fresh rollouts of the generic prompt over the same window. κ is the content analogue of α: 1 if the class evidence is kept, 0 if the rollout looks like a fresh generic rollout.

Drift metrics for prompt-only rollouts. Where no opening is held, drift is measured against fresh 5-s generations of the same prompt, in colour statistics (D) and DINOv2 [17] features (K; Appendix B).

Pairing and preregistration. With bf16 kernels, the same seed and code give bit-identical videos within a GPU architecture but different videos across architectures, so arms that are compared run in one job, which regenerates its outputs if a preempted job resumes on another architecture. Each study's hypotheses, decision rules and analysis code were frozen by SHA-256 hash before its videos were generated (study 1: before analysis), and studies 3–36 were committed publicly first. Tests within a study are Holm-corrected; Table 3 lists every primary test.

4 Rollouts relax toward a conditioning-set default

Tone. Observationally, the 31-s change of every global statistic is anti-correlated with its opening value (r from -0.79 to -0.95; Appendix Table 5), but openings differ in content too, so we intervened: we shifted the brightness and colour temperature of the same opening within each of 24 scenes (12 NAVSIM, 12 generated) and continued every version with the generic prompt (study 6, 31 s; study 10, 60 s). The versions converge (Fig. 1). At 60 s every cache configuration has relaxed (Appendix Fig. 5a): α ranges from -0.01 (LongLive without a sink) to 0.79 (Self Forcing GAN), ranked as at 31 s (Spearman $\rho = 0.78$, $p = 0.008$). Colour temperature behaves like brightness, the response is linear in the shift from -24 to $+24$ L* (study 23), and high persistence is not a frozen video (Appendix C). Persistence varies by checkpoint even with architecture, context and sampler fixed (Self Forcing: 0.35 to 0.79 at 60 s; two SiD checkpoints differ by 0.36), so the training objective alone does not set it.

4.1 Scene content

Tone is not what a user sees as "the scene". We generated openings of six scene classes (snow, desert, forest, coast, mountain, farmland; four each, selected by CLIP before any continuation) and continued each under the generic city prompt and under its own class prompt, with eight configurations (study 9). Content relaxes like tone (Fig. 2). Under the city prompt, κ at 30 s ranges from 0.07 (SiD) to 0.61 (Rolling Forcing), and the opening's class is CLIP's top label in only 34% of Self Forcing frames at 30 s. Under the opening's own class prompt the class persists (κ 0.71–1.05; 9.3 [8.1, 10.5] more log-odds, preregistered), as a fresh rollout of that prompt would. Across the eight configurations, content and tone persistence are correlated, but not significantly ($\rho = 0.67$, $p = 0.07$). Scoring frames whose CIELAB mean and spread are matched to the opening's, or scoring with SigLIP or DINOv2, gives the same verdicts (SigLIP: 18 of 19 tests; Appendix C).

Sinks act differently on content. Self Forcing's sink keeps more of it at 60 s (κ 0.61 vs 0.14, +0.48 [0.39, 0.57], preregistered), but not steadily: κ falls to 0.32 at 30 s and rises again, and the dunes, the sea and the mountains of the opening reappear after a minute of city (Fig. 2, left). The returns are positional (study 16, 24 openings): with the sink's keys re-rotated to sit just before the window, they vanish (burst contrast, κ over 50–55 s minus 45–50 s: -0.29 [-0.37, -0.22]), and the content kept at 60 s falls from 0.56 to 0.25 (-0.31 [-0.37, -0.25]). Three new seeds per opening (study 21) reproduce the sink's advantage, the loss under re-rotation (-0.40 [-0.47, -0.34]) and the return, and suggest a small share kept by the re-rotated sink: 0.19 against 0.11 without a sink (+0.08, Holm $p = 0.006$, though an interval that also resamples the fresh starts is [-0.01, 0.16]). LongLive's sink, which keeps tone, does not keep content (0.22 vs 0.24, preregistered test not supported).

Content relaxes outside the Wan-1.3B family too. FramePack, continuing the same 24 openings under the city prompt, keeps 0.55 of the class evidence at 20 s (Self Forcing: 0.40 with a sink, 0.31 without; study 17). Giving its two routes different images separates them (Appendix Fig. 12b): the clean context prefix alone keeps 0.24 [0.12, 0.37], levelling off after 14 s, while the SigLIP image embedding alone, with a city frame in

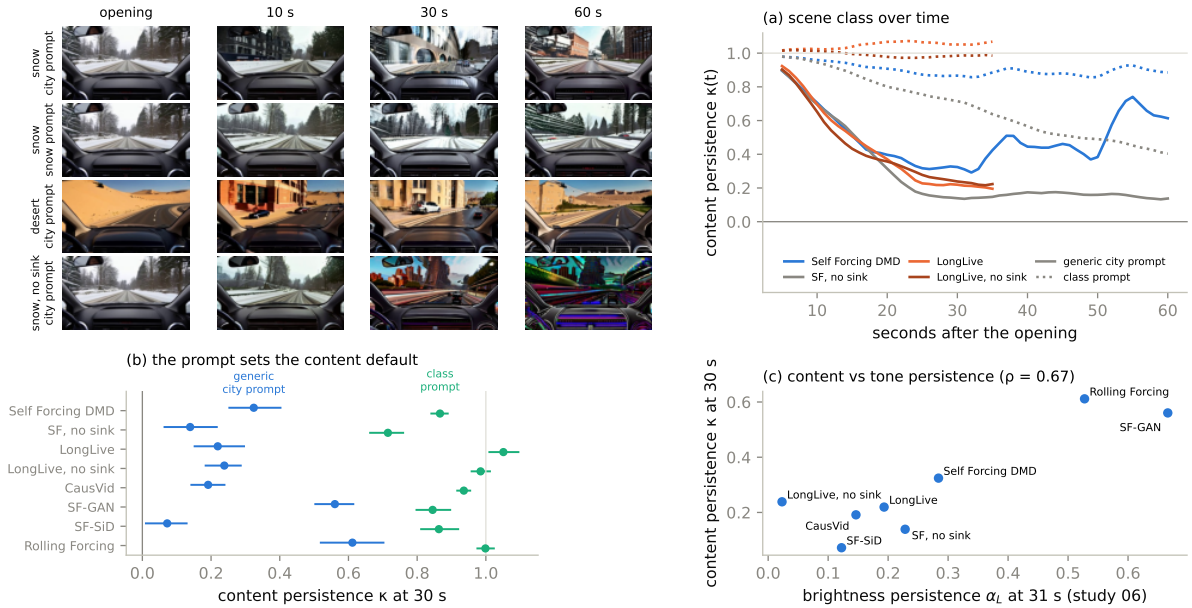


Figure 2: Scene content relaxes to the prompt's default (study 9). Left: a snow and a desert opening continued by Self Forcing DMD under the city prompt or their own class prompt. Under the city prompt both become city streets; with a sink the desert returns at 60 s; without a sink the video collapses. (a) Content persistence $\kappa(t)$ (CLIP log-odds of the opening's class against "city street", 0 = fresh generic rollout, 1 = opening) under the city prompt (solid) and the class prompt (dotted). (b) κ at 30 s, 24 openings, 95% CIs. (c) Content against brightness persistence across configurations.

the context, *pulls* the video toward the opening's class, to 0.23 [0.16, 0.30] at 20 s and still rising. The same embedding has no effect on tone (study 11), so in FramePack it sets the default of the scene but not of its brightness. Tone and content can also part the other way: Krea Realtime 14B's first-frame anchor keeps the opening's class (κ 0.37 vs 0.05 without it at 60 s, +0.33 [0.18, 0.47], study 22; the class is CLIP's top label in 37% of frames against 14%) but none of its tone (Section 5). MAGI-1 behaves like the Wan models (study 19, preregistered): under the city prompt its class score falls by 10.7 [9.4, 11.8] by 25–30 s, its own class prompt sets a different default (+7.3 [5.8, 8.8]), and full history keeps $\kappa = 0.43$ against 0.29 with the released window (+0.14 [-0.01, 0.30], not supported; +0.19 and significant on SigLIP embeddings, Appendix C).

4.2 The default

The conditioning sets it. Continuing the same 40 openings under four lighting prompts gives four defaults: on Self Forcing, sunny 43, overcast 34, dusk 19 and night 16 L*; on LongLive the same words give 45, 40, 32 and 21 (Appendix Fig. 13). Under one prompt the default also differs between checkpoints of one architecture (zero-drift levels from 17 to 54 L*), and it tracks where prompt-only rollouts end ($r = 0.93$, slope 1.58; Appendix C). Regenerating one chunk from a bit-exact copy of the cache with only its text conditioning swapped measures the prompt's pull per chunk: 0.088 [0.045, 0.149] at 15 s (study 14), as a linear recursion fitted to 120-s curves had predicted (Appendix D).

Persistence depends on the scene. Equation 1 separates a model term (α) from a prompt term (μ), but α itself varies: Self Forcing keeps 0.34 of the opening under a sunny prompt and 0.05 under a night prompt (exploratory interaction, $p = 0.001$). In four new scene types with their own prompts and self-generated openings (study 12), its persistence at 60 s differs between them (0.44 to 0.67, $p = 0.049$, preregistered; not significant across all tests), as does LongLive's at 30 s (0.10 to 0.43). The sink effect replicates almost exactly outside driving scenes: 0.55 with a sink and 0.02 without (+0.53 [0.45, 0.61]).

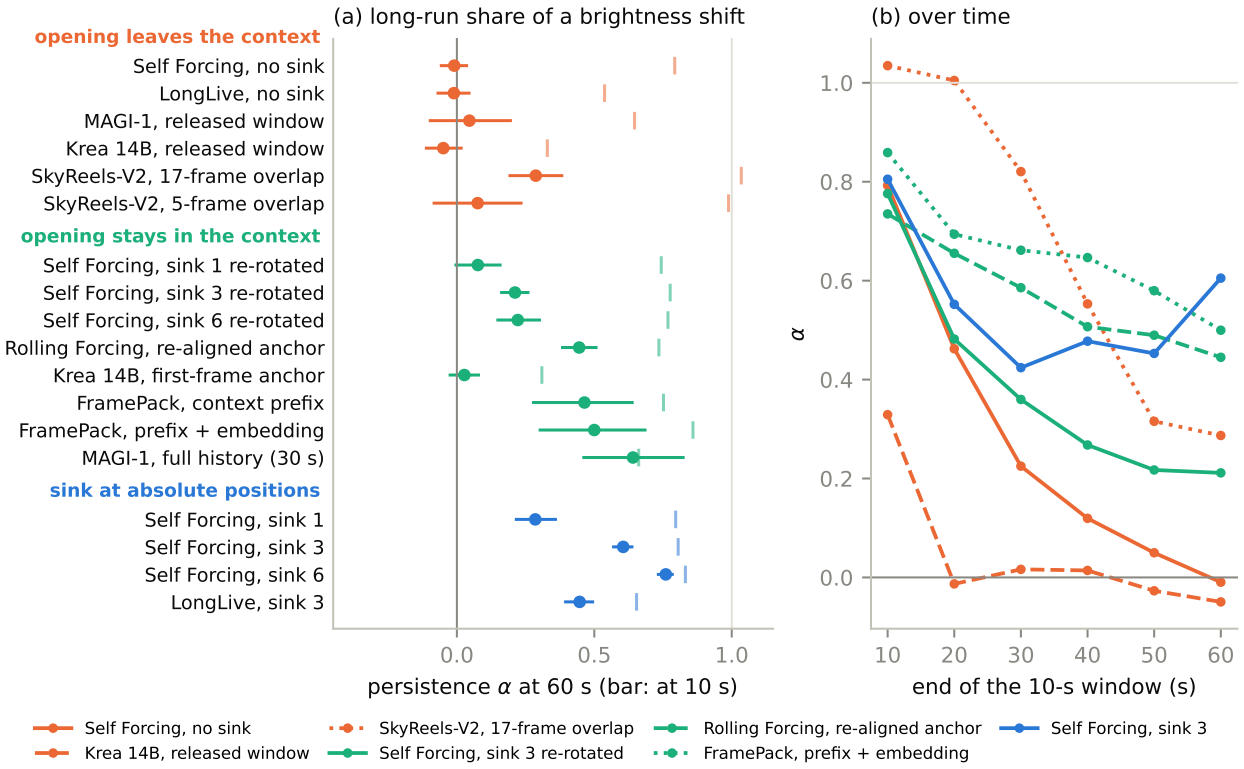


Figure 3: Whether the opening stays in the context decides how much of it is kept. (a) Persistence of an L* shift at 60 s (window [50, 60) s; MAGI-1 full history: 30 s) with scene-bootstrap 95% CIs; ticks: at 10 s. (b) Persistence over time (10-s windows) for one or two configurations per group.

Two components of drift. For a prompt-only rollout, Eq. 1 says nothing about why the default should differ from where the same prompt starts. It does, and by how much depends on the model (Appendix Fig. 10a). Over 40 prompts (study 30, preregistered), fresh 60-s rollouts of Self Forcing without a sink end 16.4 L* from where they start (mean |end−start|) and those of CausVid 11.7, against 6.8 for LongLive, tuned from Self Forcing on its own long rollouts (+9.6 [6.6, 12.4], +4.9 [2.0, 7.7]); with re-aligned sinks Self Forcing ends 9.6 away and Rolling Forcing 3.3 (+6.3 [4.2, 8.6]). Yet LongLive still forgets its opening ($\alpha = -0.01$ without a sink at 60 s). Drift thus has two separable parts: *forgetting* the opening, $(1-\alpha)(x_{\text{open}} - \mu)$, which every attention-based model does, and an *attractor offset* $\mu - x_{\text{start}}$. The offset is set less by training than we expected: it grows with the key–value window up to 21 frames (Appendix Fig. 11; studies 33 and 36, preregistered) in Self Forcing (5.2 L* at 6 frames, 16.4 at 21; +0.63 per frame [0.49, 0.77]) and LongLive (+2.1 [0.6, 3.5], 12 to 21), not in CausVid (−0.5 [-1.9, 0.9]); beyond 21 it did not grow significantly (at 24 frames within ±2 L* in both); with LongLive's window Self Forcing is not significantly further than LongLive (study 32). A sink at its absolute positions, which keeps returning to the opening, hides these differences (3.8 vs 3.0 L*). Fresh rollouts of Krea 14B brighten by about 4 L* in 30 s, those of MAGI-1 by 15–20. Prompt-only drift metrics mix the two parts (Appendix H); relaxation describes the memory of the opening, not the absence of degradation.

5 What decides how much of the opening is kept

Figure 3 collects every configuration we ran for 60 s, grouped by whether the opening stays in the model's context. Once it leaves, it is forgotten; while it stays, a share is kept whose size depends on the model; and sinks left at their absolute rotary positions add bursts in which the opening returns.

Once the opening leaves the context, it is forgotten. Without a sink, Self Forcing and LongLive keep 0.05 and −0.01 at 60 s (study 10). MAGI-1 [31] (4.5B, distilled; another architecture) attends in its released configuration to the four previous 1-s chunks, so the opening leaves its context after about 4 s; persistence falls from 0.87 at 5 s to 0.34 at 30 s (−0.52 [-0.70, -0.35], study 18), and to 0.05 [-0.10, 0.20] at 60 s (study 19, 12 scenes; from 0.21 at 30 s, a fall of −0.17 [-0.41, 0.08] that is not significant). Krea Realtime 14B [32], a Self Forcing student of Wan2.1-14B, recomputes its cache over a window of three latent frames (study 22, 24 scenes): it keeps 0.33 at 10 s and nothing from 20 s on (−0.05 at 60 s; −0.38 [-0.45, -0.31] from 10 to 60 s), and its class evidence falls to $\kappa = 0.05$. The pattern holds at 14B, though Krea's window also differs. Re-denoising overlapping frames delays forgetting but does not stop it. SkyReels-V2-DF, which re-denoises the last 17 frames of each 97-frame segment with the next, holds the opening for 20 s ($\alpha = 0.96$ [0.86, 1.07], study 6; 1.00 in study 16) and then forgets, to 0.29 at 60 s (study 16; with a 5-frame overlap, 0.08; Appendix E).

An opening that stays in the context keeps a share of its tone, its scene or both. Self Forcing keeps the sink's keys at their absolute rotary positions, ever further from the current chunk. Re-rotated to sit just before the window, as in Rolling Forcing and Deep Forcing, the sink stays at query distances seen in training. Re-rotated sinks of three and six frames keep the same share at 60 s (0.23 and 0.22, −0.00 [-0.11, 0.10], study 16; the preregistered depth test is not supported) and one frame keeps 0.08 (study 10); with three new seeds per scene a three-frame sink keeps 0.21 against −0.01 without one (+0.22 [0.15, 0.29], study 20). Rolling Forcing, which re-aligns its anchor and was trained on longer rollouts, keeps 0.45. Krea's optional first-frame anchor, one latent frame next to its window, keeps the scene (Section 4.1) but not tone: $\alpha = 0.03$ [-0.03, 0.08] (both tone tests not supported), close to the 0.08 of a one-frame re-rotated sink in Self Forcing. FramePack [23] (13B, fine-tuned from HunyuanVideo [24]) places the opening as a clean prefix frame in the context of *every* 33-frame section and also feeds its SigLIP embedding as a global condition. The shift travels through the prefix, not the embedding (0.01 [-0.14, 0.17] through the embedding alone, study 11); at 60 s FramePack keeps 0.51 [0.29, 0.71] (0.46 with the prefix alone, study 16), partly in videos that stopped moving (Appendix C). MAGI-1 with full history keeps 0.64 [0.45, 0.83] at 30 s (+0.30 [0.17, 0.43] over the released window, study 18). These shares are still falling at 60 s. At 180 s (study 24, preregistered) the re-rotated sink still keeps 0.12 [0.04, 0.20] and Rolling Forcing 0.29 [0.21, 0.37]; on new seeds to 240 s, over 60-s windows (study 29, preregistered), Rolling Forcing's share is steady from the second minute (+0.01 [-0.04, 0.06], equivalent within ±0.10), the re-rotated sink keeps 0.10 [0.06, 0.14] in the fourth minute but is not shown steady (0.18, 0.18, 0.10), and without a sink the opening stays forgotten (0.00 [-0.05, 0.05]). Keeping the opening does not stop the drift the openings share: MAGI-1's fresh rollouts brighten by 15–20 L^* under either setting, so full history preserves the difference between openings, not their level.

Sinks at their absolute positions: bursts. Without re-rotation a sink keeps more, and depth helps: no sink, one, three and six frames keep 0.05, 0.29, 0.61 and 0.76 at 60 s (+0.15 [0.11, 0.20] for six vs three; studies 7, 10), and LongLive's sink 0.45 against −0.01. Much of this is positional. Persistence jumps near 17, 33 and 51 s in every configuration with a sink of three or more frames at its absolute positions, across five Self Forcing checkpoints and LongLive (Appendix Fig. 5b, c), never with a re-rotated sink or without one, and the opening's scene returns at these times (Section 4.1). Re-rotation lowers persistence at 60 s from 0.61 to 0.16 (−0.44 [-0.54, -0.35], preregistered), and by −0.27 [-0.38, -0.16] between the bursts (40–50 s). An attention probe locates the bursts (exploratory; Fig. 4a): the share of attention a chunk's queries would give the sink if its keys sat at each query–sink distance peaks sharply at 71, 134, 203 and 266 latent frames (0.11–0.17 against a median of 0.07), identically in every chunk and opening, so the peaks are positional, not content; a sink at its original positions passes the first three 16.5, 32.3 and 49.5 s after the opening, the times of the bursts. Moving the sink moves the bursts (study 25, preregistered; Fig. 4b): 18 or 36 latent frames further back, the probe puts every burst 4.5 or 9 s earlier, and persistence rose in the predicted rather than the original windows (+0.53 [0.45, 0.61] and +0.25 [0.21, 0.30]). A sink held at one distance does not burst, and the probe's share at that distance sets how much is kept (Fig. 4c): 0.86 at 60 s at the peak distance 203 and 0.01 at the trough before it (+0.85 [0.78, 0.91], study 25), and over 14 further distances chosen by a fixed rule, persistence rose with the share (mean within-scene Spearman $\rho = +0.83$ [0.79, 0.86], study 27, preregistered), also among the 7 lowest shares (+0.35 [0.20, 0.49]). The price is motion ($\rho = -0.81$ [-0.84, -0.79]): every distance with a share of 0.11 or more froze the video on its opening (Appendix Figs. 6, 7), every distance below kept moving, and the change was abrupt (exploratory). So does LongLive (study 28, preregistered):

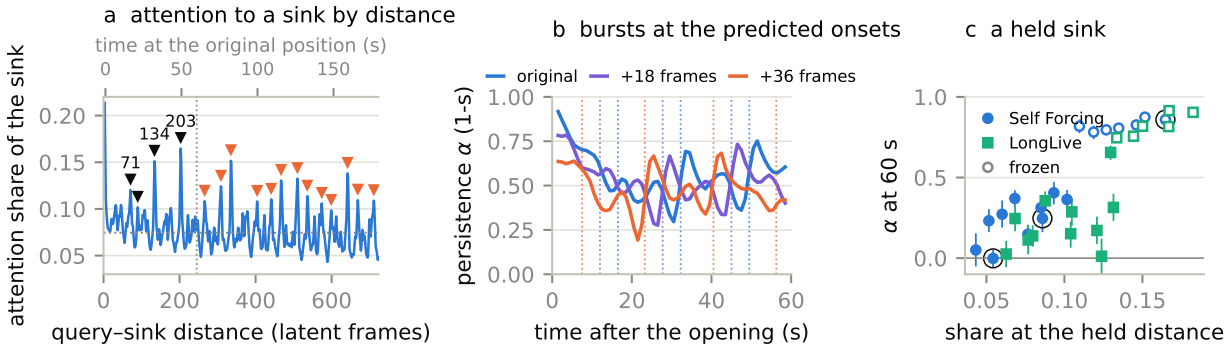


Figure 4: Bursts happen where the sink draws attention. (a) Share of attention that Self Forcing's queries would give to a three-frame sink at each query–sink distance (exploratory probe; 4 openings, 6 chunks each, all layers and passes). Triangles: major peaks (prominence ≥ 0.04); orange: the 13 after 60 s, preregistered as burst predictions (study 26). Top axis: when a sink at its original positions is at that distance. (b) Persistence of an L^* shift in 1-s windows (3-s moving average, 24 scenes) with the sink at its original positions and moved 18 or 36 frames further back (study 25); dotted lines: burst onsets predicted from the peaks in (a). (c) Persistence at 60 s with the sink held at one query–sink distance, against the probe's share there: 14 distances chosen by a rule fixed in advance (study 27) and the three of study 25 (circled); squares: LongLive against its own probe (study 28); hollow: most videos frozen on the opening.

moving its sink moves its bursts (+0.43 and +0.36), and held, its persistence rises with its own probe's share ($\rho = +0.81$ [0.78, 0.84]), correlating 0.94 with Self Forcing's across distances. A burst is thus a brief copy of the opening while the sink's distance passes a peak. A second probe predicted 13 further onsets to 180 s, and persistence jumped at them (study 26, preregistered; +0.19 [0.17, 0.20] in Self Forcing, +0.13 [0.12, 0.15] in LongLive; Appendix Fig. 8b). The peaks come from one of the 22 temporal rotary frequencies (exploratory band probe; Appendix Fig. 9): without the one of period 34 latent frames no peak remains, and holding only it at the re-rotated distance in generation removed the bursts and cut persistence at 60 s to 0.16, in Self Forcing (from 0.58; study 34) and LongLive (from 0.40; burst contrast +0.30 [0.23, 0.37]; study 35), while holding its neighbour of period 51 did not remove them (both preregistered; two finer predictions failed).

A world model. In Matrix-Game-2.0 [50], driven by actions without text, a shift held only in its 6-frame context fades (−0.24 [-0.35, -0.14] from 5 to 30 s, study 15, preregistered), while one sent through its CLIP image embedding persists (0.25 at 15 s, exploratory; Appendix E).

6 Two levers

Equation 1 suggests two ways to keep a long rollout near its opening without retraining: raise α, or move μ toward the opening.

Raising α: a content-aware correction. At denoising step 1 of 4, the noisy chunk is denoised once more with a temporary cache that holds only the anchor (the opening chunk). The difference between the per-channel spatial means of the two predictions (the DC tone) is added to the real prediction, scaled by β; the main cache and the random stream are untouched. In a preregistered confirmation on 40 held-out prompts × 4 seeds (study 2), colour drift D fell by 1.48 [1.19, 1.76] at $\beta = 1$ and by 1.78 [1.45, 2.11] at $\beta = 1.5$, with structure drift K non-inferior, and in the relaxation study the correction raised Self Forcing's persistence from 0.28 to 0.40 (+0.12 [0.05, 0.20], preregistered). Table 1 compares it with two baselines. Post-hoc colour matching to the anchor reaches a similar D but, by construction, removes legitimate global changes too: on 7 held-out prompts with real brightness transitions (tunnels, garages, clouds), matching keeps 7–12% of the transition and the correction a median of 83% (per-prompt values from 0.58 to 1.51, so on some prompts it amplifies the change). A deeper attention sink, as in Deep Forcing [9] (study 13, bit-identical pairing), is a strong alternative: nine sink frames cut D by more than the correction at $\beta = 1$ (−0.45 [-0.75, -0.15]) with

Table 1: Correction versus post-hoc colour matching and deeper sinks (study 2 confirmation set: 40 held-out prompts $\times$ 4 seeds; transitions: 7 held-out prompts with legitimate brightness changes). Transitions kept: median over prompts of the corrected over the uncorrected first-to-last brightness change, prompt-bootstrap 95% CI.

Method	ΔD [95% CI]	transitions kept	jitter	boundary jumps
Content-aware correction, $\beta = 1$	-1.48 [-1.76, -1.19]	0.83 [0.67, 1.47]	$\times 1.13$	$\times 1.32$
Content-aware correction, $\beta = 1.5$	-1.78 [-2.11, -1.45]	0.93 [0.83, 1.48]	$\times 1.67$	$\times 2.40$
Post-hoc mean matching	-0.89 [-1.29, -0.50]	0.07	$\times 0.15$	$\times 0.12$
Post-hoc Reinhard [20]	-1.58 [-1.96, -1.19]	0.12	$\times 0.13$	$\times 0.10$
Deeper sink, 6 frames (study 13)	-1.30 [-1.68, -0.93]	0.70 [0.58, 1.12]	$\times 1.04$	$\times 0.98$
Deeper sink, 9 frames (study 13)	-1.93 [-2.31, -1.55]	0.59 [0.19, 1.09]	$\times 1.16$	$\times 1.09$

smaller boundary jumps, but a sink pulls everything toward the opening, and in two of the seven transition prompts it erased most of a legitimate change (7% and 19% kept, against 151% and 147% for the correction). The correction costs one extra forward pass per chunk (about 25% more compute) and more frame-to-frame jitter; it is the more conservative choice when real changes matter.

Raising α without bursts: placing the sink. The attention probe also says where to put a sink. Held at distance 66 (below the freezing threshold, chosen from study 27 on one seed), Self Forcing's three-frame sink kept 0.45 [0.40, 0.49] of the opening in the third minute on a new seed (study 31, preregistered): twice the re-rotated sink's 0.22 (+0.23 [0.15, 0.31]) and as much as the sink at its original positions (0.43), without bursts (equivalent within ± 0.05) and with more motion than the original positions (+1.10 [0.69, 1.57] grey levels), as much as the re-rotated sink. Distance 448 kept less (+0.07 [-0.00, 0.15] over re-rotated, not supported), and in LongLive distance 66 keeps only 0.17 (study 28), so the distance has to be chosen per model.

Moving μ: prompts. Since the prompt sets the default, lighting words can be chosen so that the default matches the opening. Choosing the lighting sentence whose *measured* default is nearest the opening's brightness reduces $|\Delta L|$ at 31 s by 26% on Self Forcing (-2.45 [-4.08, -1.04], Holm $p = 0.013$, study 4). A free-form Qwen2.5-VL [18] caption of the opening does not reduce brightness drift (study 3), and the other five preregistered prompt tests failed, including LongLive with its own calibration (study 5). The mapping from words to brightness is model-specific; Appendix F shows how well Eq. 1 forecasts these outcomes.

7 Discussion

A finite-memory sampler that is ergodic under fixed conditioning forgets its initial state, so relaxation is expected; we measure how much is kept and what decides it. Across nine models the answer is whether the opening stays in the context and, for a sink at its absolute positions, where it sits: the larger shares and returning scenes of such sinks are brief copies of the opening at the distances where the sink is most attended.

A high α is not always better. Under a conflicting prompt (a snowy opening, a city prompt), forgetting the opening is following the prompt, and an anchor that keeps the opening also passes its errors forward (SkyReels amplifies the shift in 5 of 12 scenes). Persistence is better treated as a property to choose and to report at a stated horizon, with the default and its offset. A world model conditioned on a frame and a description drifts toward the description's default. In Self Forcing the mechanism that keeps the opening's scene is the unstable, positional one: keeping the opening's content but not its errors is the open problem.

Limitations. Seven of the nine models are Wan2.1-based (FramePack and MAGI-1 are not). Content relaxation was replicated on FramePack and MAGI-1, and its prompt dependence on MAGI-1, but the decomposition rests on Wan-based models, and the world model collapses after 20–30 s. Interventions shift global tone or select a scene class (CLIP and DINOv2 scores are coarse proxies for content); openings are held images; scene content is compared across openings rather than shifted within one; most configurations use one seed per scene (the main Self Forcing results replicate with more); with Holm correction across all 125 primary tests, 67 of the 91 supported ones remain significant; the correction was not compared with Pathwise or FreqForcing.

References

[1] T. Yin, Q. Zhang, R. Zhang, W. T. Freeman, F. Durand, E. Shechtman, X. Huang. From slow bidirectional to fast autoregressive video diffusion models. CVPR 2025 (arXiv:2412.07772).

[2] X. Huang, Z. Li, G. He, M. Zhou, E. Shechtman. Self Forcing: Bridging the train-test gap in autoregressive video diffusion. arXiv:2506.08009, 2025.

[3] B. Chen, D. Martí Monsó, Y. Du, M. Simchowitz, R. Tedrake, V. Sitzmann. Diffusion Forcing: Next-token prediction meets full-sequence diffusion. NeurIPS 2024 (arXiv:2407.01392).

[4] Wan Team. Wan: Open and advanced large-scale video generative models. arXiv:2503.20314, 2025.

[5] T. Yin, M. Gharbi, R. Zhang, E. Shechtman, F. Durand, W. T. Freeman, T. Park. One-step diffusion with distribution matching distillation. CVPR 2024 (arXiv:2311.18828).

[6] M. Zhou, H. Zheng, Z. Wang, M. Yin, H. Huang. Score identity distillation: Exponentially fast distillation of pretrained diffusion models for one-step generation. ICML 2024 (arXiv:2404.04057).

[7] S. Yang, W. Huang, R. Chu, Y. Xiao, Y. Zhao, X. Wang, M. Li, E. Xie, Y. Chen, Y. Lu, S. Han, Y. Chen. LongLive: Real-time interactive long video generation. arXiv:2509.22622, 2025.

[8] K. Liu, W. Hu, J. Xu, Y. Shan, S. Lu. Rolling Forcing: Autoregressive long video diffusion in real time. arXiv:2509.25161, 2025.

[9] J. Yi, W. Jang, P. H. Cho, J. Nam, H. Yoon, S. Kim. Deep Forcing: Training-free long video generation with deep sink and participative compression. arXiv:2512.05081, 2025.

[10] SkyReels Team. SkyReels-V2: Infinite-length film generative model. arXiv:2504.13074, 2025.

[11] G. Xiao, Y. Tian, B. Chen, S. Han, M. Lewis. Efficient streaming language models with attention sinks. ICLR 2024 (arXiv:2309.17453).

[12] X. Xiang, Z. Duan, G. Zhang, H. Zhang, Z. Gao, J. Wu, S. Zhang, T. Wang, Q. Fan, C. Guo. Pathwise test-time correction for autoregressive long video generation. arXiv:2602.05871, 2026.

[13] J. Li, L. Liang, L. Kong, Y. Zhang. FreqForcing: Autoregressive long video generation via spectral self-anchoring. arXiv:2607.27110, 2026.

[14] J. Wang, F. Zhang, X. Li, V. Y. F. Tan, T. Pang, C. Du, A. Sun, Z. Yang. Error analyses of auto-regressive video diffusion models: A unified framework. arXiv:2503.10704, 2025.

[15] Z. Huang et al. VBench: Comprehensive benchmark suite for video generative models. CVPR 2024 (arXiv:2311.17982).

[16] J. Ke, Q. Wang, Y. Wang, P. Milanfar, F. Yang. MUSIQ: Multi-scale image quality transformer. ICCV 2021.

[17] M. Oquab et al. DINOv2: Learning robust visual features without supervision. TMLR 2024 (arXiv:2304.07193).

[18] Qwen Team. Qwen2.5-VL technical report. arXiv:2502.13923, 2025.

[19] D. Dauner et al. NAVSIM: Data-driven non-reactive autonomous vehicle simulation and benchmarking. NeurIPS 2024.

[20] E. Reinhard, M. Ashikhmin, B. Gooch, P. Shirley. Color transfer between images. IEEE CG&A, 2001.

[21] J. Cui, J. Wu, M. Li, T. Yang, X. Li, R. Wang, A. Bai, Y. Ban, C.-J. Hsieh. LoL: Longer than longer, scaling video generation to hour. arXiv:2601.16914, 2026.

[22] J. Cui, J. Wu, M. Li, T. Yang, X. Li, R. Wang, A. Bai, Y. Ban, C.-J. Hsieh. Self-Forcing++: Towards minute-scale high-quality video generation. arXiv:2510.02283, 2025.

[23] L. Zhang, S. Cai, M. Li, G. Wetzstein, M. Agrawala. Frame context packing and drift prevention in next-frame-prediction video diffusion models. arXiv:2504.12626, 2025.

[24] W. Kong et al. HunyuanVideo: A systematic framework for large video generative models. arXiv:2412.03603, 2024.

[25] J. Kim, J. Kang, J. Choi, B. Han. FIFO-Diffusion: Generating infinite videos from text without training. NeurIPS 2024.

[26] R. Henschel et al. StreamingT2V: Consistent, dynamic, and extendable long video generation from text. CVPR 2025 (arXiv:2403.14773).

[27] H. Qiu, M. Xia, Y. Zhang, Y. He, X. Wang, Y. Shan, Z. Liu. FreeNoise: Tuning-free longer video diffusion via noise rescheduling. ICLR 2024.

[28] Y. Lu, Y. Liang, L. Zhu, Y. Yang. FreeLong: Training-free long video generation with SpectralBlend temporal attention. NeurIPS 2024.

[29] D. Ruhe, J. Heek, T. Salimans, E. Hoogeboom. Rolling diffusion models. ICML 2024.

[30] Y. Jin et al. Pyramidal flow matching for efficient video generative modeling. ICLR 2025.

[31] Sand AI. MAGI-1: Autoregressive video generation at scale. arXiv:2505.13211, 2025.

[32] Krea AI. Krea Realtime 14B: real-time, long-form AI video generation. Technical blog post and model release, https://www.krea.ai/blog/krea-realtime-14b, 2025.

[33] K. Song, B. Chen, M. Simchowitz, Y. Du, R. Tedrake, V. Sitzmann. History-guided video diffusion. ICML 2025 (arXiv:2502.06764).

[34] S. Bengio, O. Vinyals, N. Jaitly, N. Shazeer. Scheduled sampling for sequence prediction with recurrent neural networks. NeurIPS 2015.

[35] S. Ross, G. Gordon, D. Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. AISTATS 2011.

[36] M. Ning, E. Sangineto, A. Porrello, S. Calderara, R. Cucchiara. Input perturbation reduces exposure bias in diffusion models. ICML 2023.

[37] M. Ning, M. Li, J. Su, A. A. Salah, I. O. Ertugrul. Elucidating the exposure bias in diffusion models. ICLR 2024.

[38] D. Valevski, Y. Leviathan, M. Arar, S. Fruchter. Diffusion models are real-time game engines. ICLR 2025.

[39] E. Alonso, A. Jelley, V. Micheli, A. Kanervisto, A. Storkey, T. Pearce, F. Fleuret. Diffusion for world modeling: Visual details matter in Atari. NeurIPS 2024.

[40] S. Gao, J. Yang, L. Chen, K. Chitta, Y. Qiu, A. Geiger, J. Zhang, H. Li. Vista: A generalizable driving world model with high fidelity and versatile controllability. NeurIPS 2024.

[41] S. Lin, B. Liu, J. Li, X. Yang. Common diffusion noise schedules and sample steps are flawed. WACV 2024.

[42] N. Guttenberg. Diffusion with offset noise. CrossLabs blog, 2023.

[43] X. Gu, T. Pang, C. Du, Q. Liu, F. Zhang, C. Du, Y. Wang, M. Lin. When attention sink emerges in language models: An empirical view. ICLR 2025.

[44] T. Darcet, M. Oquab, J. Mairal, P. Bojanowski. Vision transformers need registers. ICLR 2024.

[45] K. Li, T. Liu, N. Bashkansky, D. Bau, F. Viégas, H. Pfister, M. Wattenberg. Measuring and controlling instruction (in)stability in language model dialogs. COLM 2024.

[46] P. D. Oldham. A note on the analysis of repeated measurements of the same subjects. Journal of Chronic Diseases 15:969–977, 1962.

[47] A. G. Barnett, J. C. van der Pols, A. J. Dobson. Regression to the mean: what it is and how to deal with it. International Journal of Epidemiology 34(1):215–220, 2005.

[48] A. Radford et al. Learning transferable visual models from natural language supervision. ICML 2021.

[49] X. Zhai, B. Mustafa, A. Kolesnikov, L. Beyer. Sigmoid loss for language image pre-training. ICCV 2023.

[50] X. He, C. Peng, Z. Liu, B. Wang et al. Matrix-Game 2.0: An open-source, real-time, and streaming interactive world model. arXiv:2508.13009, 2025.

[51] M. Zhao, G. He, Y. Chen, H. Zhu, C. Li, J. Zhu. RIFLEx: A free lunch for length extrapolation in video diffusion transformers. ICML 2025 (arXiv:2502.15894).

[52] F. Barbero, A. Vitvitskyi, C. Perivolaropoulos, R. Pascanu, P. Veličković. Round and round we go! What makes rotary positional encodings useful? arXiv:2410.06205, 2024.

A Models, preregistered tests, deviations and withdrawn claims

Table 2: Configurations. Wan-based cache models generate 832×480 at 16 fps in chunks of 3 latent frames (0.75 s); window and sink sizes are in latent frames.

Configuration	Base; training	Steps	Context of the opening
Self Forcing DMD, SiD, SiD v2, GAN [2, 5, 6]	Wan2.1-1.3B; self-rollout on 5-s clips	4	window 21; sink 0/1/3/6, original or re-rotated positions
Self Forcing 10 s	Wan2.1-1.3B; self-rollout on 10-s clips	4	window 21, sink 3
CausVid [1]	Wan2.1-1.3B; teacher-forced DMD	3	window 21, sink 0/3
LongLive [7]	Wan2.1-1.3B; long-rollout tuning	4	window 12, sink 0/3
Rolling Forcing [8]	Wan2.1-1.3B; rolling joint denoising	5	anchor block + 21
Krea Realtime 14B [32]	Wan2.1-14B; Self Forcing distillation	4	key–value cache recomputed from the last 3 latent frames; optional first-frame anchor (1 of the 3)
SkyReels-V2-DF [10]	Wan2.1-1.3B architecture; diffusion forcing; 544p, 24 fps	30 + CFG	17 (or 5) re-denoised overlapping frames
Matrix-Game-2.0 [50]	Wan2.1-1.3B architecture, no text; actions; 352×640, 25 fps	3	first latent frame only, window 6, no sink; CLIP image embedding
FramePack-F1 / I2V [23]	HunyuanVideo 13B; 480p, 30 fps	25	clean prefix in every section; SigLIP image embedding
MAGI-1 4.5B-distill [31]	own architecture (chunk-wise autoregressive diffusion transformer); 432×768, 24 fps, chunks of 1 s	64, pipelined	released: window of 4 earlier chunks, no sink; or full history

Table 3: All preregistered primary tests and their outcomes. Intervals are Wald intervals from the preregistered mixed models in studies 1, 2 and 13 and bootstrap intervals otherwise; p is Holm-adjusted within each study; *supported but not significant after Holm correction across all 125 tests. Study 3 LongLive rows are as executed; after regenerating four cross-architecture openings the conclusions are unchanged; in study 1, 62% of nominal pairs crossed GPU architectures, which widens but does not bias its intervals. Units: D, K drift metrics; $|\Delta L|$ and |off| (end minus start of fresh rollouts) in L*; α persistence, κ content persistence. Arms: G generic prompt, M VLM caption of the opening, C calibrated lighting plus caption, ttc content-aware correction. ‡The t-test treats the default μ as fixed; the bootstrap interval also resamples the fresh starts that estimate it (pooled with seed 0: $[0.02, 0.15]$). §Equivalence by two one-sided tests; p is the larger one-sided p.

Study	Test	Model	Estimate [95% CI]	p_{Holm}	Outcome
1 Context	sink 3 vs none, D	SF	-7.56 [-8.32, -6.81]	$< 10^{-16}$	supported
	sink 3 vs none, K	SF	-0.90 [-1.04, -0.76]	$< 10^{-16}$	supported
	window × sink, D	SF	$+5.71$ [4.21, 7.22]	$< 10^{-12}$	supported
	window × sink, K	SF	$+0.67$ [0.38, 0.95]	$< 10^{-5}$	supported
2 Correction	ttc β=1, D	SF	-1.48 [-1.76, -1.19]	$< 10^{-23}$	supported
	ttc β=1.5, D	SF	-1.78 [-2.11, -1.45]	$< 10^{-24}$	supported
3 Caption	M − G, $\lvert\Delta L\rvert$	SF	-1.21 [-3.09, 0.36]	0.58	not supported
	M+ttc − M, $\lvert\Delta L\rvert$	SF	-1.00 [-1.89, -0.12]	0.14	not supported
	M − G, $\lvert\Delta L\rvert$	LL	-0.19 [-2.00, 1.63]	1.0	not supported
	M+ttc − M, $\lvert\Delta L\rvert$	LL	-0.36 [-1.48, 0.75]	1.0	not supported
4 Calibrated	C − M, $\lvert\Delta L\rvert$	SF	-1.24 [-2.54, 0.12]	0.25	not supported
	C − G, $\lvert\Delta L\rvert$	SF	-2.45 [-4.08, -1.04]	0.013	supported*
	C − M, $\lvert\Delta L\rvert$	LL	-1.15 [-2.59, 0.29]	0.26	not supported
	C − G, $\lvert\Delta L\rvert$	LL	-1.34 [-3.19, 0.55]	0.26	not supported
5 Calibrated (LL)	C − M, $\lvert\Delta L\rvert$	LL	-1.36 [-3.06, 0.32]	0.26	not supported
	C − G, $\lvert\Delta L\rvert$	LL	-1.55 [-3.50, 0.39]	0.26	not supported
6 Relaxation	no sink − sink, α_L (31 s)	SF	-0.06 [-0.17, 0.05]	0.34	not supported
	no sink − sink, α_L (31 s)	LL	-0.17 [-0.28, -0.06]	0.013	supported*
	ttc − none, α_L (31 s)	SF	$+0.12$ [0.05, 0.20]	0.013	supported*

Study	Test	Model	Estimate [95% CI]	p_{Holm}	Outcome
	no sink − sink, α_b (31 s)	SF	−0.26 [-0.34, -0.18]	$< 10^{-5}$	supported
7 Sink at 60 s	sink − no sink, α_L	SF	+0.55 [0.45, 0.65]	$< 10^{-9}$	supported
	sink − no sink, α_L	CausVid	+0.19 [0.14, 0.24]	$< 10^{-6}$	supported
8 Overlap	5 − 17 frames, α_L (20 s)	SkyReels	−0.24 [-0.40, -0.09]	0.014	supported*
9 Content	class − city prompt, class score (30 s)	SF	+9.33 [8.11, 10.51]	$< 10^{-9}$	supported
	sink − no sink, κ (30 s)	LL	−0.02 [-0.10, 0.07]	0.72	not supported
	sink − no sink, κ (60 s)	SF	+0.48 [0.39, 0.57]	$< 10^{-7}$	supported
10 Atlas	sink 6 − sink 3, α_L (60 s)	SF	+0.15 [0.11, 0.20]	$< 10^{-6}$	supported
	sink 1 − no sink, α_L (60 s)	SF	+0.24 [0.14, 0.34]	< 0.0002	supported
	sink − no sink, α_L (60 s)	LL	+0.46 [0.39, 0.53]	$< 10^{-10}$	supported
	re-rotated − original sink, α_L (60 s)	SF	−0.44 [-0.54, -0.35]	$< 10^{-7}$	supported
11 FramePack	α (both routes, 20 s) > 0.5	FramePack-F1	0.68 [0.57, 0.79]	0.033	supported*
	α (embedding only, 20 s) > 0	FramePack-F1	0.01 [-0.14, 0.17]	0.87	not supported
	both − context only, α (20 s)	FramePack-F1	+0.14 [-0.05, 0.33]	0.38	not supported
12 Prompts	sink − no sink, α_L (60 s), 4 domains	SF	+0.53 [0.45, 0.61]	$< 10^{-10}$	supported
	sink − no sink, α_L (30 s), 4 domains	LL	+0.19 [0.11, 0.28]	0.005	supported*
	α_L differs between domains (ANOVA)	SF	$F = 3.12$	0.049	supported*
13 Deep sink	sink 6, D	SF	−1.30 [-1.68, -0.93]	$< 10^{-6}$	supported
	sink 9, D	SF	−1.93 [-2.31, -1.55]	$< 10^{-6}$	supported
14 Prompt pull	$w_p > 0$ at 15 s (replay)	SF	+0.088 [0.045, 0.149]	0.006	supported*
15 World model	ctx-only shift fades (30 vs 5 s)	Matrix-Game	−0.24 [-0.35, -0.14]	0.024	supported*
	emb-only shift persists, α (30 s)	Matrix-Game	+0.43 [0.13, 0.71]	0.075	not supported
	both − ctx-only, α (30 s)	Matrix-Game	+0.33 [-0.07, 0.71]	0.18	not supported
16 Re-rotation	re-rotated − original, burst contrast R	SF	−0.29 [-0.37, -0.22]	$< 10^{-5}$	supported
	re-rotated − original sink, κ (60 s)	SF	−0.31 [-0.37, -0.25]	$< 10^{-6}$	supported
	re-rotated sink − none, κ (60 s)	SF	+0.10 [-0.00, 0.21]	0.11	not supported
	re-rotated − original sink 6, α_L (60 s)	SF	−0.54 [-0.62, -0.45]	$< 10^{-9}$	supported
	re-rotated sink 6 − sink 3, α_L (60 s)	SF	−0.00 [-0.11, 0.10]	0.94	not supported
	α (both routes, 60 s) > 0	FramePack-F1	0.51 [0.29, 0.71]	0.004	supported*
	α (context only, 60 s) > 0	FramePack-F1	0.46 [0.28, 0.65]	0.003	supported*
	5 frames: α (60 s) − α (20 s)	SkyReels	−0.57 [-0.79, -0.39]	0.001	supported*
	17 frames: α (60 s) − α (20 s)	SkyReels	−0.72 [-0.81, -0.63]	$< 10^{-6}$	supported
17 FP content	content falls, class score (20 s)	FramePack-F1	−6.72 [-8.29, -5.20]	$< 10^{-5}$	supported
	κ (context only, 20 s) > 0	FramePack-F1	0.24 [0.12, 0.37]	0.001	supported*
	κ (embedding only, 20 s) > 0	FramePack-F1	0.23 [0.16, 0.30]	$< 10^{-6}$	supported
	both − context only, κ (20 s)	FramePack-F1	+0.31 [0.20, 0.41]	$< 10^{-4}$	supported
18 MAGI-1	window: α (30 s) − α (5 s)	MAGI-1	−0.52 [-0.70, -0.35]	< 0.0004	supported
	full history − window, α (30 s)	MAGI-1	+0.30 [0.17, 0.43]	0.002	supported*
	α (full history, 30 s) > 0	MAGI-1	0.64 [0.45, 0.83]	< 0.0002	supported
19 MAGI-1	content falls, class score (25–30 s)	MAGI-1	−10.66 [-11.82, -9.45]	$< 10^{-8}$	supported
	full history − window, κ (25–30 s)	MAGI-1	+0.14 [-0.01, 0.30]	0.10	not supported
	class − city prompt, class score (25–30 s)	MAGI-1	+7.28 [5.84, 8.78]	$< 10^{-6}$	supported
	window: α (60 s) − α (30 s)	MAGI-1	−0.17 [-0.41, 0.08]	0.22	not supported
20 Seeds	sink 3 − none, α (60 s)	SF	+0.62 [0.56, 0.67]	$< 10^{-15}$	supported
	re-rotated − original sink, α (60 s)	SF	−0.39 [-0.46, -0.32]	$< 10^{-8}$	supported
	re-rotated sink − none, α (60 s)	SF	+0.22 [0.15, 0.29]	$< 10^{-5}$	supported
21 Content seeds	sink 3 − none, κ (60 s)	SF	+0.48 [0.40, 0.56]	$< 10^{-10}$	supported
	re-rotated − original sink, κ (60 s)	SF	−0.40 [-0.47, -0.34]	$< 10^{-12}$	supported
	re-rotated sink − none, κ (60 s)	SF	+0.08 [-0.01, 0.16]‡	0.006	supported*
	re-rotated − original, burst contrast R	SF	−0.34 [-0.38, -0.28]	$< 10^{-10}$	supported
22 Krea	window: α (60 s) − α (10 s)	Krea 14B	−0.38 [-0.45, -0.31]	$< 10^{-9}$	supported
	anchor − window, α (60 s)	Krea 14B	+0.08 [-0.03, 0.18]	0.32	not supported
	α (anchor, 60 s) > 0	Krea 14B	0.03 [-0.03, 0.08]	0.37	not supported
	content falls, class score (55–60 s)	Krea 14B	−12.02 [-12.64, -11.37]	$< 10^{-15}$	supported
	anchor − window, κ (55–60 s)	Krea 14B	+0.33 [0.18, 0.47]	$< 10^{-5}$	supported
23 Dose–response	sink: α ($\pm$24) − α ($\pm$8), $\pm$0.15§	SF	+0.03 [-0.03, 0.10]	0.005	supported*
	re-rotated sink: the same§	SF	−0.01 [-0.14, 0.11]	0.020	supported*
	sink: brightening − darkening within $\pm$0.15§	SF	+0.01 [-0.06, 0.08]	0.003	supported*
	no sink: α ($\pm$24) within $\pm$0.10§	SF	+0.05 [0.01, 0.08]	0.005	supported*
24 180-s horizon	re-rotated sink, α (180 s) > 0	SF	+0.12 [0.04, 0.20]	0.044	supported*
	Rolling Forcing, α (180 s) > 0	RF	+0.29 [0.21, 0.37]	$< 10^{-5}$	supported
	re-rotated sink steady, 60→180 s, $\pm$0.10§	SF	−0.09 [-0.22, 0.03]	0.55	not supported
	Rolling Forcing steady, 60→180 s, $\pm$0.10§	RF	−0.07 [-0.16, 0.01]	0.55	not supported
	no sink, α (180 s) within $\pm$0.10§	SF	+0.04 [-0.06, 0.14]	0.38	not supported
25 Sink position	sink 18 frames back: bursts move	SF	+0.53 [0.45, 0.61]	$< 10^{-10}$	supported
	sink 36 frames back: bursts move	SF	+0.25 [0.21, 0.30]	$< 10^{-9}$	supported
	held at peak (203) − trough (187), α (60 s)	SF	+0.85 [0.78, 0.91]	$< 10^{-16}$	supported
	held at peak (203) − distance 20, α (60 s)	SF	+0.64 [0.54, 0.74]	$< 10^{-10}$	supported
26 Burst times	bursts at 13 predicted onsets, 65–177 s	SF	+0.19 [0.17, 0.20]	$< 10^{-16}$	supported
	bursts at the same onsets	LL	+0.13 [0.12, 0.15]	$< 10^{-13}$	supported
	predicted − control onsets	SF	+0.32 [0.29, 0.35]	$< 10^{-15}$	supported
27 Held distance	persistence rises with the share (ρ)	SF	+0.83 [0.79, 0.86]	$< 10^{-16}$	supported
	motion falls with the share (ρ)	SF	−0.81 [-0.84, -0.79]	$< 10^{-16}$	supported
	the same among the 7 lowest shares (ρ)	SF	+0.35 [0.20, 0.49]	0.0002	supported

Study	Test	Model	Estimate [95% CI]	p_{Holm}	Outcome
28 LL sink position	sink 18 frames back: bursts move	LL	$+0.43$ $[0.36, 0.50]$	$< 10^{-10}$	supported
	sink 36 frames back: bursts move	LL	$+0.36$ $[0.29, 0.43]$	$< 10^{-8}$	supported
	held: persistence rises with LL's share (ρ)	LL	$+0.81$ $[0.78, 0.84]$	$< 10^{-16}$	supported
	held: motion falls with the share (ρ)	LL	-0.83 [-0.85, -0.81]	$< 10^{-16}$	supported
29 Steady share	re-rotated sink, α (180–240 s) > 0	SF	$+0.10$ $[0.06, 0.14]$	0.0005	supported
	Rolling Forcing, α (180–240 s) > 0	RF	$+0.32$ $[0.27, 0.38]$	$< 10^{-9}$	supported
	re-rotated sink steady, 2nd→4th min, $\pm 0.10^{\S}$	SF	-0.08 [-0.13, -0.02]	0.20	not supported
	Rolling Forcing steady, the same§	RF	$+0.01$ [-0.04, 0.06]	0.002	supported*
	no sink, α (60–240 s) within $\pm 0.10^{\S}$	SF	-0.00 [-0.05, 0.05]	0.002	supported*
30 Attractor offset	|off|: SF − LL, no sink	SF, LL	$+9.59$ $[6.62, 12.45]$	$< 10^{-6}$	supported
	|off|: CausVid − LL, no sink	CausVid, LL	$+4.89$ $[1.98, 7.67]$	0.002	supported*
	|off|: SF re-rotated sink − RF	SF, RF	$+6.33$ $[4.22, 8.61]$	$< 10^{-5}$	supported
31 Held-sink lever	held at 66 − re-rotated, α (120–180 s)	SF	$+0.23$ $[0.15, 0.31]$	$< 10^{-4}$	supported
	held at 448 − re-rotated, α (120–180 s)	SF	$+0.07$ [-0.00, 0.15]	0.074	not supported
	held at 66 − original positions, motion	SF	$+1.10$ $[0.69, 1.57]$	0.0002	supported
	held at 66: burst contrast within $\pm 0.05^{\S}$	SF	-0.02 [-0.03, -0.01]	$< 10^{-5}$	supported
32 Window-matched	|off|: SF window 12 − LL, no sink	SF, LL	$+1.27$ [-0.97, 3.56]	0.28	not supported
	|off|: SF window 12 − LL, sink 3	SF, LL	$+1.08$ [-0.31, 2.40]	0.27	not supported
33 Window dose	|off| slope over windows 6–21 (per frame)	SF	$+0.63$ $[0.49, 0.77]$	$< 10^{-9}$	supported
	|off|: LongLive window 21 − 12	LongLive	$+2.07$ $[0.64, 3.47]$	0.013	supported*
	|off|: CausVid window 21 − 12	CausVid	-0.52 [-1.91, 0.93]	0.48	not supported
34 Band hold	sink frequency j=4 held: bursts go	SF	$+0.33$ $[0.26, 0.39]$	$< 10^{-8}$	supported
	j=8–11 held: bursts stay, $\pm 0.10^{\S}$	SF	-0.06 [-0.11, -0.00]	0.13	not supported
	j=3 held: bursts at the probe's new times	SF	$+0.05$ [0.00, 0.11]	0.13	not supported
	j=4 held: less kept, α (50–60 s)	SF	$+0.42$ $[0.35, 0.49]$	$< 10^{-9}$	supported
35 Band hold II	sink frequency j=4 held: bursts go	LL	$+0.30$ $[0.23, 0.37]$	$< 10^{-6}$	supported
	j=4 held: less kept, α (50–60 s)	LL	$+0.24$ $[0.16, 0.33]$	$< 10^{-4}$	supported
	j=5 held − j=4 held, bursts	SF	$+0.30$ $[0.24, 0.37]$	$< 10^{-7}$	supported
	j=5 held − j=4 held, bursts	LL	$+0.23$ $[0.16, 0.30]$	$< 10^{-5}$	supported
36 Window beyond	|off|: SF window 24 − 21	SF	-0.23 [-1.52, 1.08]	1	not supported
	|off|: SF window 27 − 21	SF	$+1.86$ $[0.30, 3.43]$	0.081	not supported
	|off|: LongLive window 24 − 21	LongLive	$+0.08$ [-0.74, 0.87]	1	not supported

Deviations from the preregistrations.

- *Study 1* was frozen before analysis but after generation of its main videos had started, and its arms were not pinned to one GPU architecture: 62% of its nominal seed pairs ran on different architectures. Its preregistered Wald intervals are reported; the prompt-cluster bootstrap interval for the main contrast is wider ([-10.45, -5.39] for D) but leads to the same conclusion.
- *Study 3:* a preempted job resumed on another GPU type and split four LongLive openings across architectures. We regenerated the affected arms on the original architecture; conclusions are unchanged.
- *Study 6:* the ODE-initialised Self Forcing student was replaced by SiD v2 before generation (it collapses to blur after the opening). GPU pinning was replaced by an architecture guard that regenerates a task's outputs when a resumed job lands on another architecture; paired arms of each preregistered contrast always share a task. Exploratory configurations (sink 1; correction on LongLive and CausVid) were added after freezing but before any data, from separate task lists, so their contrasts with the preregistered arms are not guaranteed to share an architecture. Prompt-only reference rollouts were renamed after generation; no analysis code changed.
- *Study 7* reuses the 24 scenes of study 6, including the 6 whose 120-s curves motivated it. Restricted to the 18 other scenes, the sink effect is $+0.53$ $[0.42, 0.63]$ (Self Forcing) and $+0.18$ $[0.13, 0.24]$ (CausVid).
- *Studies 10 and 11* were run on two non-preemptible 4-GPU jobs and A40-only array tasks instead of the preregistered all-queue arrays, because no large GPU was free on the shared queue; all their videos ran on A40 GPUs, so pairing is unaffected.
- *Study 16:* its last two arms were preregistered for A40 GPUs only; with no A40 free, half of the 12 scenes were moved to the Ada group (L40S, RTX 6000 Ada) before any of their videos existed, keeping both versions of a scene on one architecture. Its 60-s SkyReels videos ran out of GPU memory in the final post-processing step, which a wrapper moved to the CPU without changing the frozen generator (differences within the 0.2 L^* run-to-run variation of SkyReels).
- *Study 22:* Krea's released code does not output the starting frame, so the frozen generator held the opening for 8 frames instead of 9. One line was corrected after 32 tone videos had been generated; for

those, the per-frame statistics were completed with a copy of frame 0, which is identical to what the corrected generator writes. The analysis code was not changed. The other studies ran as preregistered.

Withdrawn or revised claims. During the project we withdrew four claims after checks: that MUSIQ rewards collapsed videos (it is merely insensitive); that the sink stabilises the default (trend difference +0.5 [-2.4, 3.5]); that the model follows its context by 76% (true only for artificial shifts); and that early dynamics forecast the final brightness. The recursion's predicted Self Forcing plateau of 0.25–0.54 was missed (0.58 at 60 s in study 7).

B Drift metrics for prompt-only rollouts

For rollouts without a held opening, drift is measured against 10 (Self Forcing) or 5 (LongLive) fresh 5-s generations of the same prompt. D is the per-frame Mahalanobis distance of (Lab means, saturation, log sharpness) from the fresh frames' distribution, and K the mean cosine distance of a frame's DINOv2 feature to its 5 nearest fresh-frame features, normalised by the fresh clips' own value; both are averaged over seconds 29–31. Transitions kept is the median over held-out transition prompts of the corrected over the uncorrected first-to-last brightness change; values above 1 mean the correction amplified the change.

C Robustness

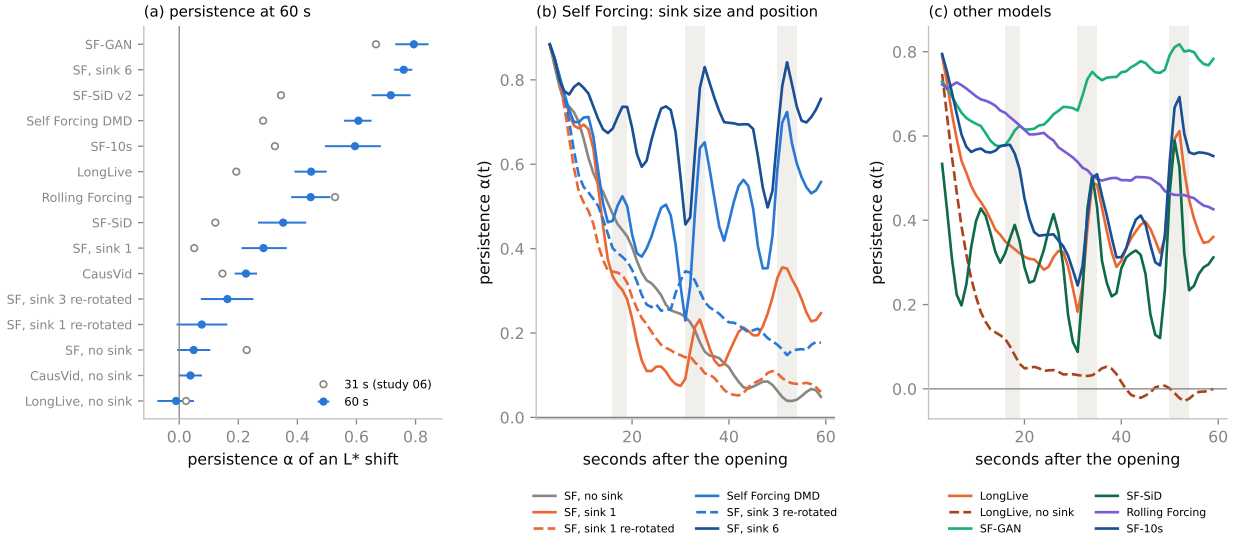


Figure 5: Persistence at 60 s (study 10; CausVid from study 7). (a) α of an L* shift with scene-bootstrap 95% CIs; hollow markers: 31 s (study 6). (b) Self Forcing with sinks of 0, 1, 3 and 6 latent frames, at their original positions or re-rotated to sit just before the window (dashed). (c) Other models. Curves: 1-s windows, 3-s moving average; shaded: 16–19 s, 31–35 s and 50–54 s, where sinks of three or more frames at their original positions produce bursts. The 60-s endpoint window, [50, 60) s, overlaps the third burst; Appendix C reports a burst-free window.

Immediate loss. Part of the loss is immediate: most configurations lose a fifth to a third of the shift within 4.5 s, and SiD loses two thirds there and nothing afterwards, which looks like rejecting an off-manifold shift rather than forgetting it.



Figure 6: Frames of two scenes (+16 L* opening) continued by Self Forcing with a three-frame sink at its original positions and held at query–sink distances 203 (a probe peak), 187 (the trough before it) and 20 (next to the window), study 25. Held at the peak, the video stays on its opening; at the trough it leaves the opening within seconds, with streak artefacts. Left: persistence α and mean frame-to-frame change (grey levels) over 50–60 s, 24 scenes.

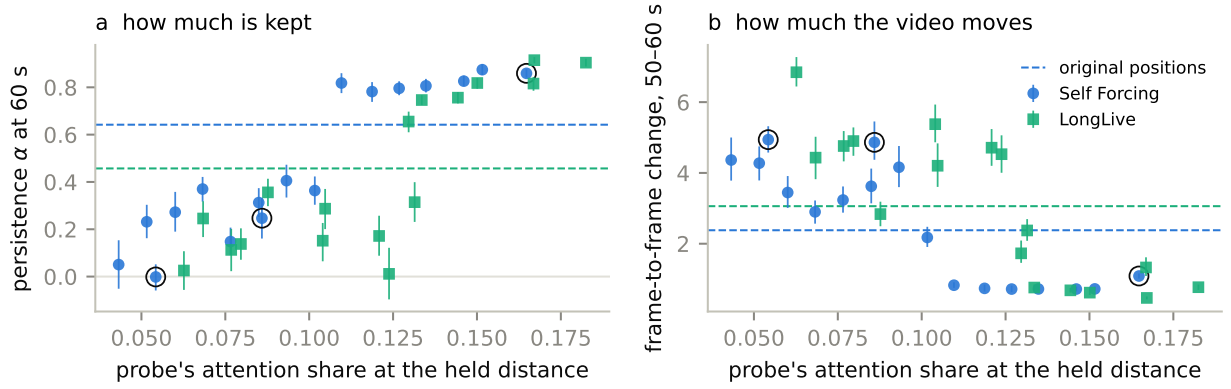


Figure 7: The sink held at one query–sink distance, 60-s videos, 24 scenes: Self Forcing (study 27; circled: the three distances of study 25) and LongLive against its own probe's shares (study 28). (a) Persistence α and (b) mean frame-to-frame change over 50–60 s against the probe's attention share at that distance, with scene-bootstrap 95% CIs. Dashed: the sink at its original positions (blue Self Forcing, green LongLive). Above a share of about 0.105 (Self Forcing) or 0.13 (LongLive) almost every distance freezes the video on its opening; below it the video moves and persistence rises with the share, more noisily. Frozen distances keep 0.78–0.87 with a frame-to-frame change of 0.7–1.1 grey levels, against 2.4 with the sink in place.

Linearity. Every α in the main text comes from shifts of ± 16 L*. Shifting the same 24 openings by -24 to $+24$ L* (study 23, Self Forcing, 60 s, 720 videos), α for ± 24 and ± 8 agree with the sink in place (0.61 and 0.58) and re-rotated (0.20 and 0.21), darkening and brightening are kept alike (0.60 and 0.61), and without a sink the large shift is forgotten too (0.05 $[0.01, 0.08]$); the four equivalence tests held within the study but not after correction across all tests. The least-squares line of the end brightness on the opening's over five shifts (-24 to $+24$ L*) has slope 0.60 $[0.56, 0.64]$ with the sink in place, 0.20 $[0.16, 0.24]$ re-rotated and 0.05 $[0.01, 0.08]$ without a sink, and no quadratic term ($p = 0.88$, 0.16, 0.31). The burst at 31–34 s scales with the shift (jump 0.40 for ± 8, 0.34 for ± 24), so a burst copies the opening itself. A re-rotated sink keeps a darkening better than a brightening (0.26 vs 0.15; not tested). In study 6, for 12 of 14 configurations the persistence estimated from ± 8 and from ± 16 shifts differs by at most 0.09 at 31 s, with intervals including 0. SiD and the sink-1 configuration give near-zero slopes for small shifts (-0.10, -0.15) and positive ones for large shifts (0.18, 0.10); LongLive with a sink keeps a darkening better than a brightening (0.36 vs 0.04). The linear law is a first-order description.

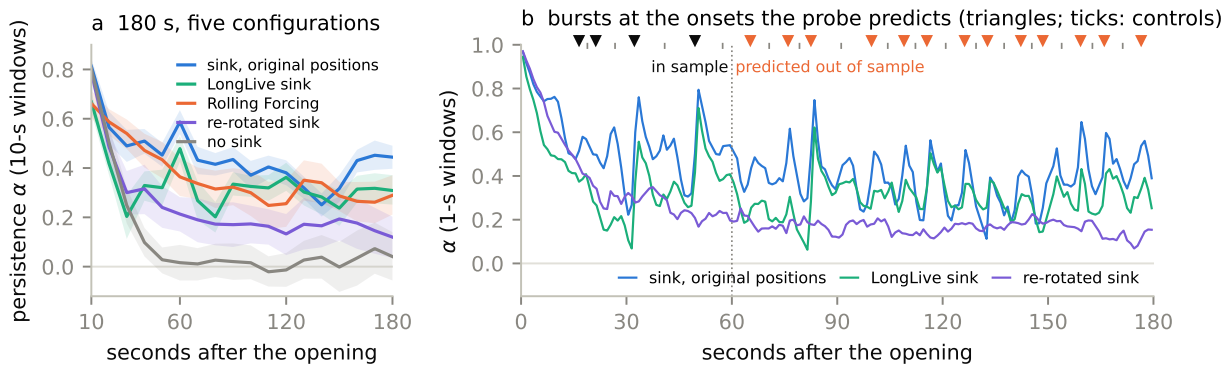


Figure 8: The 180-s horizon (studies 24 and 26; 24 scenes, Self Forcing unless named). (a) Persistence over 10-s windows with scene-bootstrap 95% CIs. An opening that stays in the context keeps a share to 180 s; without a sink it is forgotten by 50 s. (b) Scene-mean 1-s curves. Triangles: burst onsets of the probe's major peaks for a sink at its original positions; the four before 60 s were seen when the probe was made, the 13 after 60 s were preregistered predictions (study 26). Ticks: control onsets midway between them. The re-rotated sink does not burst.

The 60-s endpoint and the positional bursts. The preregistered 60-s window, 50–60 s, overlaps the third burst. Over 40–50 s, Self Forcing keeps 0.48 with a three-frame sink and 0.20 with a re-rotated one (difference −0.27 [-0.38, -0.16]); over the full period 35–60 s, 0.53 and 0.19. Sink depth is monotone in every window (40–50 s: no sink 0.08, one frame 0.23, three 0.48, six 0.63). Some curves are not monotone decays (GAN and SiD v2 lose part of the shift early and regain it), and Rolling Forcing is still declining at 60 s. To 180 s (study 24, Fig. 8a), the shares of Rolling Forcing and the re-rotated sink are the same in the second and third minutes over 60-s windows (0.30 and 0.30; 0.17 and 0.16; exploratory), but the preregistered equivalence tests over 10-s windows failed; on new seeds to 240 s (study 29) Rolling Forcing's share passed the 60-s-window test and the re-rotated sink's did not; a sink at its original positions keeps 0.44 at 180 s, +0.32 [0.21, 0.45] over the re-rotated one. At the 13 onsets after 60 s that the 180-s probe predicted (study 26), Self Forcing's jump was +0.32 [0.29, 0.35] larger than at control onsets between them, and 12 of the 13 large jumps after 60 s followed a predicted onset (chance 0.33). In the 120-s subset of study 6 (6 scenes), Self Forcing with a sink keeps 0.54, 0.48 and 0.44 at 60, 90 and 118 s; on this subset a preregistered replication of study 26 found bursts at the 6 predicted onsets between 60 and 120 s (Self Forcing +0.23 [0.16, 0.31], LongLive +0.20 [0.11, 0.31]; predicted − control onsets +0.32 [0.21, 0.45]). No value we report is shown to be an asymptote; Eq. 1 is a statement at a stated horizon T, and α and μ are always reported with theirs.

Which rotary frequencies make the bursts. Wan2.1's temporal rotary embedding rotates 22 frequency pairs with periods $2\pi \cdot 10000^{2j/44}$ latent frames ($j = 0, \ldots, 21$: 6.3, 9.5, 14.5, 22.1, 33.5, 51.0, …, $4 \cdot 10^4$). A band probe (exploratory; 2 openings, 2 chunks) placed the sink's keys at distance 20 and rotated one band of frequencies to each distance of the grid. Without $j = 4$ (period 33.5) no major peak remains (peak contrast 0.058 to 0.018; Fig. 9a); $j = 4$–7 alone give a comb with a tooth every ≈34 frames (69, 103, 169, 203, 237), and $j = 3$ (period 22.1 ≈ 33.5/1.5) cancels every other tooth: without it the odd teeth return (33, 103, 165, 234). Removing $j = 8$–11, $j = 12$–21 or any single j of 5–11 leaves the curve unchanged ($r \geq 0.99$); $j = 0$–2 shape the peaks. The peaks are present in all 30 layers, in 272 of 360 layer–head pairs and at every denoising step. Study 34 (preregistered; 24 scenes, a new seed, 60 s) held these bands of the sink at distance 20 in generation, the others at their original positions (Fig. 9b). Holding $j = 4$ removed the bursts (burst contrast at the original onsets −0.02 against +0.31; difference +0.33 [0.26, 0.39]) and lowered persistence at 50–60 s from 0.58 to 0.16 (+0.42 [0.35, 0.49]), below the fully re-rotated sink's 0.28 (−0.11 [-0.19, -0.03], exploratory), while the video kept moving (motion 4.0 against 2.3; no frozen video). Holding $j = 8$–11 kept persistence (0.59) and the bursts at 32 and 50 s but made them smaller (−0.06 [-0.11, -0.00]), so equivalence within ±0.10 was not shown. Holding $j = 3$ raised the contrast at the four new onsets the probe predicts only slightly (+0.05 [0.00, 0.11], not supported; +0.10 at 24.5 s): the probe's finer prediction did not carry over to generation. With this seed the sink at its original positions did not burst at 16.5 s (−0.02), so whether holding

Table 4: Persistence over the same window, 15–20 s after the opening, for every model except Krea and MAGI-1 (exploratory; scene-bootstrap 95% CIs; Matrix-Game: scenes without a black second before 20 s).

Configuration	α	95% CI	scenes
SkyReels-V2-DF, 17 overlap frames	1.01	[0.89, 1.15]	12
FramePack-I2V, anti-drifting sampler	0.86	[0.79, 0.92]	12
SkyReels-V2-DF, 5 overlap frames	0.77	[0.64, 0.90]	12
Self Forcing DMD, sink 6	0.72	[0.65, 0.78]	24
FramePack-F1, both routes	0.68	[0.56, 0.78]	12
Rolling Forcing	0.64	[0.60, 0.67]	24
Self Forcing GAN	0.61	[0.54, 0.67]	24
Self Forcing 10-s	0.54	[0.46, 0.63]	24
FramePack-F1, context prefix only	0.53	[0.31, 0.76]	12
Self Forcing DMD, sink 3	0.49	[0.44, 0.53]	24
Self Forcing DMD, no sink	0.44	[0.38, 0.51]	24
Self Forcing DMD, sink 3 re-rotated	0.38	[0.28, 0.46]	24
Matrix-Game-2.0, both routes	0.37	[0.22, 0.52]	17
Self Forcing SiD	0.34	[0.24, 0.46]	24
LongLive, sink 3	0.33	[0.25, 0.40]	24
Self Forcing DMD, sink 1 re-rotated	0.33	[0.25, 0.41]	24
Self Forcing DMD, sink 1	0.29	[0.17, 0.41]	24
Self Forcing SiD v2	0.24	[0.11, 0.36]	24
Matrix-Game-2.0, embedding only	0.22	[0.06, 0.36]	17
CausVid, sink 3	0.21	[0.16, 0.26]	24
Matrix-Game-2.0, context only	0.10	[-0.02, 0.22]	17
LongLive, no sink	0.08	[0.04, 0.12]	24
CausVid, no sink	0.08	[0.04, 0.12]	24
FramePack-F1, embedding only	0.01	[-0.14, 0.17]	12

Table 5: Every global statistic relaxes toward a default (observational; 40 held openings, generic prompt, 31 s). Persistence = 1 + slope of the change on the opening value; r = correlation of the change with the opening value.

	Self Forcing		LongLive	
Statistic	persistence	r	persistence	r
L^*	0.31	−0.91	0.26	−0.81
a^*	0.38	−0.92	0.39	−0.90
b^*	0.22	−0.88	0.42	−0.80
saturation	0.54	−0.92	0.23	−0.95
log sharpness	0.37	−0.91	0.23	−0.94
L^* contrast	0.34	−0.82	0.37	−0.79

$j = 3$ removes that burst could not be tested. Study 35 (preregistered; 24 scenes, a new seed) repeated the $j = 4$ hold in LongLive and Self Forcing against holding the neighbouring $j = 5$ (period 51.0). Holding $j = 4$ removed the bursts in LongLive (−0.01 against +0.29; +0.30 [0.23, 0.37]) and lowered persistence at 50–60 s from 0.40 to 0.16, and replicated in Self Forcing (+0.25 [0.19, 0.31]; 0.59 to 0.16). Holding $j = 5$ kept the bursts ($j = 5$ against $j = 4$: +0.30 [0.24, 0.37] in Self Forcing, +0.23 [0.16, 0.30] in LongLive) but is not inert: LongLive's bursts were smaller (−0.07, 90% CI [−0.12, −0.02]) and both models kept 0.06–0.10 less at 60 s, so equivalence to the original positions was not shown. Holding $j = 4$ lowered persistence already in the first 10 s, before the first burst (0.53 against 0.68 in LongLive, 0.65 against 0.79 in Self Forcing), so the frequency also carries some of what is kept between bursts.

Re-rotated sinks across studies. A re-rotated three-frame sink keeps 0.16 at 60 s in study 10 and 0.23 in study 16 because 9 of the 24 scenes ran on another GPU architecture in study 16, which amounts to resampling them; all 72 videos repeated on the same architecture were bit-identical.

Which denoising step to correct. Correcting at step 0 removes as much drift as at step 1 but erases tunnels, and steps 2–3 do nothing, consistent with global tone being fixed at the earliest denoising steps [41, 42].

Persistence is not freezing. A video frozen on its opening would score high. We recomputed motion (mean frame-to-frame grey-level change on 4×-subsampled frames; $r = 0.9999$ with the value from raw frames) from the 2272 saved videos (exploratory). Across 38 configuration–study pairs at 60 s, persistence

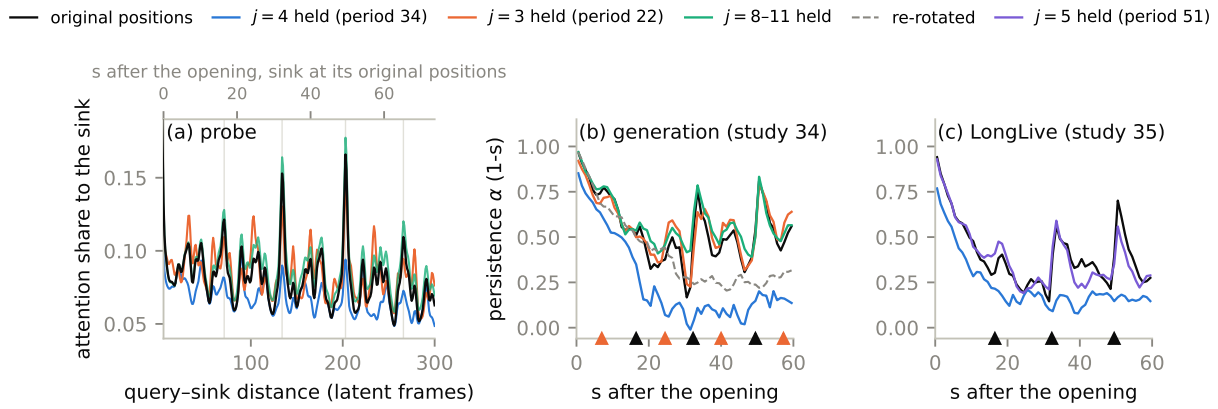


Figure 9: One rotary frequency of the sink makes the bursts. (a) Band probe (exploratory): the share of attention to a three-frame sink at each query–sink distance with all temporal frequencies rotated to that distance (black) and with one band held at distance 20 (colours); grey lines: the four major peaks. Top axis: when a sink at its original positions reaches the distance. (b) Generation with the same bands of the sink held at distance 20 (study 34; scene means of 1-s persistence, 24 scenes). Black triangles: burst onsets of the full probe; orange: the new onsets of the probe without $j = 3$. (c) LongLive with $j = 4$ or the neighbouring $j = 5$ held (study 35; a new seed).

and motion are uncorrelated (Spearman $\rho = -0.13$, $p = 0.43$), and freezing (motion below 1.5 over 50–60 s) is more common without a sink (42–58% of Self Forcing videos, 71% of CausVid's) than with one (18–29%). Restricted to scenes whose two versions both move, every Wan-based configuration keeps within 0.07 of its α over all scenes, as recomputed here (three-frame sink 0.59 vs 0.61, six frames 0.75 vs 0.76, GAN 0.73 vs 0.80, Rolling Forcing 0.44 vs 0.44); in FramePack, MAGI-1 and Matrix-Game (other frame rates) it is 0.08–0.20 lower (FramePack's prefix alone 0.26 on 6 moving scenes vs 0.46), so part of their persistence comes from videos that stopped moving. Only a sink held at a high-share distance freezes Self Forcing or LongLive on its opening (Fig. 7).

Real and generated openings. Half of the 24 tone openings are NAVSIM camera frames and half the model's own generations. At 60 s the difference in α between them averages -0.02 over 23 configuration–study pairs and no 95% interval excludes zero (at 30 s, near a burst, 4 of 23 do, in both directions). In the drift regressions of the early studies, real and generated openings fall on the same line on Self Forcing ($p = 0.80$); on LongLive, generated openings end 4.3 L* brighter than real ones of equal brightness ($p = 0.03$).

The default across horizons. The zero-drift level $\mu = b_0/(1 - \alpha)$ of the same regressions changes by at most 4.5 L* between 20 and 60 s without a sink, with re-rotated sinks and for LongLive and CausVid, but by 6–12 L* with a Self Forcing sink at its absolute positions, where α moves with the bursts and μ is ill-conditioned as α nears 1; SiD's and Rolling Forcing's rise as their rollouts brighten (by 15 and 6 L*).

Content beyond CLIP. With DINOv2 features (exploratory; class score = mean cosine similarity to the other three openings of the class minus that to the 20 openings of other classes, normalised as κ), content persistence at 30 s agrees with CLIP's across the 16 configuration–prompt cells of study 9 (Pearson $r = 0.92$). An instance score, the similarity to the video's own opening minus that to the other openings of its class (1 = the opening, 0 = an unrelated scene of the class), separates the scene from its class. Under the opening's own class prompt, where a fresh rollout would also show the class, the class persists (DINOv2 κ 0.30–0.99) but the particular scene keeps only 0.20–0.75 at 30 s; the matching prompt adds to the instance in CausVid, LongLive and Rolling Forcing (+0.19 to +0.31) but not in the Self Forcing checkpoints. A Self Forcing sink at its absolute positions brings back the particular opening, not just its class: at 60 s under the city prompt it adds +0.33 $[0.23, 0.42]$ to the instance score.

The default and fresh rollouts. Across 10 configurations, the zero-drift level correlates $r = 0.93$ with where prompt-only rollouts end, but the slope is 1.58 and the mean absolute difference 4.9 L^* (largest for the GAN and SiD v2 checkpoints, whose zero-drift levels are extrapolated from slopes far from 0). Fresh rollouts give the direction and rough size of the default, not its exact value.

Persistence depends on the prompt. In the lighting experiment, the slope of the 31-s change on the opening differs between lighting prompts on Self Forcing (prompt × opening interaction, opening-clustered Wald test $p = 0.001$): persistence is 0.34 under sunny, 0.24 overcast, 0.17 dusk and 0.05 night prompts. On LongLive the interaction is weaker ($p = 0.053$). Dark defaults compress the range of brightness the model produces, which may explain part of this. Study 12 confirms the prompt dependence in four new scene types (Section 4.2).

Content verdicts barely depend on the encoder. Every content result uses CLIP ViT-L/14 class log-odds. We rescored all content videos with SigLIP so400m-patch14-384 [49], with its own text embeddings of the same seven labels, and reran each frozen analysis unchanged (Table 6). Every verdict but one agrees: MAGI-1's full history keeps more content than its released window by +0.14 with CLIP ($p = 0.10$, not supported) and +0.19 with SigLIP ($p = 0.009$); we keep the preregistered CLIP verdict. Other estimates in κ agree within 0.03 except Krea's anchor (+0.33 and +0.44); raw log-odds differ in scale between encoders. This includes FramePack's embedding route, whose conditioning encoder is itself a SigLIP model: CLIP and SigLIP give it the same content share (0.23 and 0.25). Class scores might also carry global colour (snow is bright, sand is yellow). On frames whose CIELAB mean and standard deviation are matched to the held opening's before encoding, the city-versus-class-prompt difference of study 9 is 9.0 [7.9, 10.2] log-odds (9.3 without matching), and the four tests of study 21 give +0.45, −0.38, +0.08 and −0.32 (+0.48, −0.40, +0.08, −0.34 without).

Table 6: Content results barely depend on the scoring encoder (exploratory). Each preregistered content test rerun with its frozen analysis on SigLIP so400m embeddings instead of CLIP ViT-L/14. Estimates in each encoder's log-odds or κ units; Holm p within each study; † not supported. The verdict agrees for 18 of 19 tests; the exception is MAGI-1's full-history advantage (19 A2), not supported with CLIP, the preregistered encoder.

		CLIP		SigLIP	
Study	Test	est.	p	est.	p
09 H1	prompt sets content: $s_{\text{class}} - s_{\text{generic}}$ (30 s)	+9.33	<0.001	+13.24	<0.001
09 H2	LongLive sink keeps content: $\Delta\kappa$ (30 s)	-0.02	0.718†	+0.00	0.920†
09 H3	Self Forcing sink keeps content: $\Delta\kappa$ (60 s)	+0.48	<0.001	+0.48	<0.001
16 A1	late return is positional: ΔR	-0.29	<0.001	-0.28	<0.001
16 A2	re-rotation lowers content: $\Delta\kappa$ (60 s)	-0.31	<0.001	-0.32	<0.001
16 A3	re-rotated sink keeps content: $\Delta\kappa$ (60 s)	+0.10	0.113†	+0.09	0.099†
17 H1	FramePack content relaxes: $s - s_{\text{open}}$	-6.72	<0.001	-10.10	<0.001
17 H2	context prefix carries content: κ	+0.24	0.001	+0.22	<0.001
17 H3	SigLIP embedding carries content: κ	+0.23	<0.001	+0.25	<0.001
17 H4	embedding adds content: $\Delta\kappa$	+0.31	<0.001	+0.30	<0.001
21 H1	sink keeps content, 3 new seeds: $\Delta\kappa$ (60 s)	+0.48	<0.001	+0.50	<0.001
21 H2	re-rotation lowers content: $\Delta\kappa$ (60 s)	-0.40	<0.001	-0.40	<0.001
21 H3	re-rotated sink keeps content: $\Delta\kappa$ (60 s)	+0.08	0.006	+0.10	<0.001
21 H4	late return is positional: ΔR	-0.34	<0.001	-0.32	<0.001
22 B1	Krea content relaxes: $s - s_{\text{open}}$ (60 s)	-12.02	<0.001	-17.04	<0.001
22 B2	Krea anchor keeps content: $\Delta\kappa$ (60 s)	+0.33	<0.001	+0.44	<0.001
19 A1	MAGI-1 content relaxes: $s - s_{\text{open}}$ (30 s)	-10.66	<0.001	-14.13	<0.001
19 A2	MAGI-1 full history keeps more: $\Delta\kappa$ (30 s)	+0.14	0.103†	+0.19	0.009
19 A3	MAGI-1 prompt sets content: $s_{\text{class}} - s_{\text{generic}}$ (30 s)	+7.28	<0.001	+9.79	<0.001

D The recursion and the replay

Model a chunk's mean statistic as $x_{k+1} = w_w\,\bar{x}_{\text{window}} + w_s\,x_{\text{sink}} + w_p\,\mu$ with $w_w + w_s + w_p = 1$, where the opening counts as a window chunk until it rolls out and μ is the prompt's default. Without a sink the opening is eventually forgotten; with one, $\alpha_\infty = w_s/(w_s + w_p)$, so a small sink weight keeps a large share if the prompt's pull is also small. Fitted to 120-s curves, the recursion gives $w_s = 0.060$ [0.021, 0.097] and $w_w = 0.86$, implying $w_p \approx 0.08$, which we preregistered as the expected value of the replay measurement (the

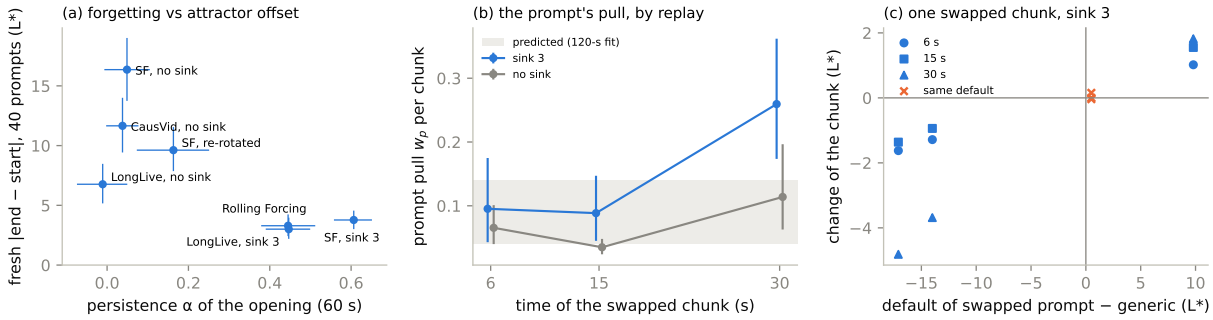


Figure 10: (a) Two components of drift: persistence of the opening at 60 s (studies 7, 10) against the attractor offset of fresh rollouts (mean |end−start| brightness over 40 prompts, 3 seeds, 60 s; study 30), 95% CIs; LongLive and Rolling Forcing have a smaller offset but still forget; Self Forcing with LongLive's 12-frame window ends 8.0 L* away (study 32, not shown). (b) The prompt's per-chunk pull measured by bit-exact replay (study 14; 16 seeds; 95% CIs): one chunk regenerated from identical cache, noise and random state with only the text conditioning swapped; grey: the value implied by the recursion fitted to 120-s curves, preregistered for 15 s. (c) The swapped chunk's change against the difference of the prompts' defaults; a prompt with the same default has no effect.

preregistered test was $w_p > 0$). Study 14 measured 0.088 $[0.045, 0.149]$ at 15 s. With a sink the prompt's pull rises to 0.26 $[0.17, 0.36]$ at 30 s (preregistered test of time invariance, +0.16 $[0.10, 0.24]$), so the recursion holds for about 20 s. Bit-exact replay of the key–value cache measures the recursion's weights directly, by regenerating one chunk from identical noise after changing only the sink, the window or the text conditioning: $w_s = 0.053$ $[0.040, 0.066]$ for synthetic tone shifts (0.036 in a dose–response design), and $w_w = 0.76$ (synthetic) to 0.90–0.99 (natural differences), in 10 development prompts × 2 seeds. The prompt's pull measured in study 14 is proportional to the difference between the two prompts' defaults, and a prompt with the same default has no effect (Fig. 10b, c). The three weights sum to 1.00 with the fitted w_w, but to 0.90 or 1.13 with the replay-measured ones, and α_∞ is 0.29–0.41 depending on which w_s is used. Without a sink the measured pull is smaller (0.035 at 15 s), so w_p is not a property of the prompt alone, and fitted weights depend on the horizon. Re-rotated sinks of three or six frames keep 0.16–0.23, below what the recursion implies; a sink at its absolute positions keeps 0.61, above it.

E Models outside the Wan cache family

SkyReels-V2-DF. It extends videos in 97-frame segments that re-denoise the last 17 frames of the previous segment jointly with the new ones. It holds the opening for 20 s ($\alpha = 0.96$ $[0.86, 1.07]$ in study 8, 1.00 in study 16) whatever its guidance, history noise or number of steps, and then forgets steadily, to 0.29 at 60 s (-0.72 [-0.81, -0.63] from 20 to 60 s, study 16); with a 5-frame overlap it relaxes earlier (-0.24 [-0.40, -0.09] at 15–20 s, preregistered) and keeps 0.08 at 60 s.

Matrix-Game-2.0. The preregistered tests of study 15 used 7 scenes, as the model often collapses to black: the context-only shift fades (-0.24 [-0.35, -0.14] from 5 to 30 s), and the embedding-only shift keeps 0.43 $[0.13, 0.71]$ at 30 s (Holm $p = 0.075$, not supported); over 20 scenes at 15 s it keeps 0.25 $[0.17, 0.36]$ (exploratory; Fig. 12d).

F The default and the forecast

How well does the law forecast? Taking α and the generic default from the relaxation study and the lighting defaults from the lighting experiment (different scenes and openings), $\Delta L \approx (1 - \alpha)(\mu - L_{\text{open}})$ predicts the 31-s brightness change of each opening in the calibrated-lighting study with $r = 0.82$ (Self Forcing) and 0.77 (LongLive), mean absolute error 3.8 L*. Most of this correlation reflects the opening's brightness. The forecast of the *benefit* of calibration per opening is weaker ($r = 0.46$ and 0.58), and its mean

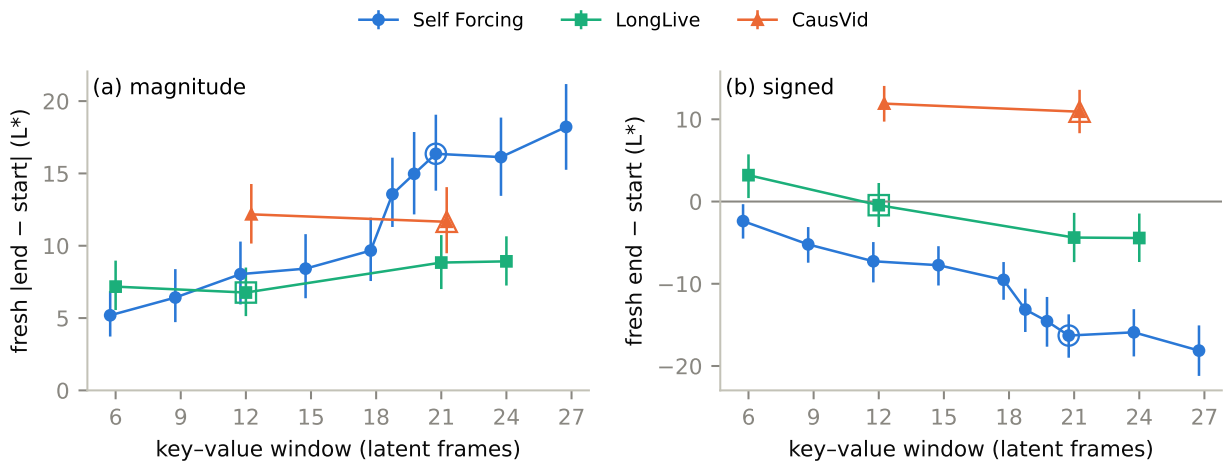


Figure 11: The attractor offset of fresh rollouts against the key–value window (studies 30, 32, 33, 36; 40 prompts, 3 seeds, no sink, [50, 60) s; prompt-bootstrap 95% CIs; circles: released windows). (a) Mean |end−start| brightness; (b) mean signed offset. In Self Forcing the offset grows with the window (preregistered slope +0.63 L* per frame [0.49, 0.77] over 6–21), most of it from 18 to 19 frames (+3.9 [2.7, 5.1]; 19 to 20: +1.4; 20 to 21: +1.4 [-0.2, 2.9]), not at the training length, and rollouts darken more; LongLive's grows from 12 to 21 frames (+2.1 [0.6, 3.5]); CausVid's does not change (−0.5 [-1.9, 0.9]). Beyond 21 frames (study 36, preregistered) neither model's offset grew significantly: at 24 frames it is within ±2 L* of 21 in both (Self Forcing −0.2 [-1.5, 1.1], LongLive +0.1 [-0.7, 0.9]); at 27 Self Forcing's is +1.9 [0.3, 3.4] larger (Holm $p = 0.08$), with more frame-to-frame change (0.77) and less sharpness (ratio 1.5 against 2.0). Short windows reach their final offset sooner (Self Forcing: 90% of it within 5 s at 6 frames, 44 s at 21) and change slightly more from frame to frame (mean $|\Delta L^*|$ 0.38 against 0.29), without losing sharpness.

is overpredicted (−4.5 predicted vs −3.2 and −2.8 observed L*). The law is a first-order guide to a prompt choice, not a precise forecast.

G Pitfalls for evaluating long-video models

1. **The GPU architecture changes the video.** With bf16, outputs are bit-identical within an architecture and differ across A40, A100/A800 and L40S/RTX 6000 Ada. Paired designs must pin or guard the architecture.
2. **Resumed jobs can silently break pairing.** Generators that skip existing outputs resume a preempted job wherever it lands.
3. **Image-quality scores miss drift.** MUSIQ, used by VBench for imaging quality, varies only between 66 and 69 across context strategies whose colour drift differs threefold.
4. **Natural variation is too small to estimate persistence.** Within a prompt, seeds start only about 2 L* apart, so within-prompt slopes of end on start brightness have intervals of about ±0.4; controlled shifts are needed.
5. **Change-on-baseline regressions mislead.** Regressing change on a baseline measured with error produces spurious regression to the mean [46, 47]. Held openings are measured almost without error, and the interventional design removes the issue.

H Context strategies (study 1)

A preregistered factorial study of 10 window and sink strategies on Self Forcing (40 prompts × 4 seeds, 31 s) found that a 3-frame sink cuts colour drift D by 7.56 [6.81, 8.32] and DINOv2 drift K by 0.90; that window size matters only without a sink; and that re-rotating the sink's positions (which keeps the query–sink

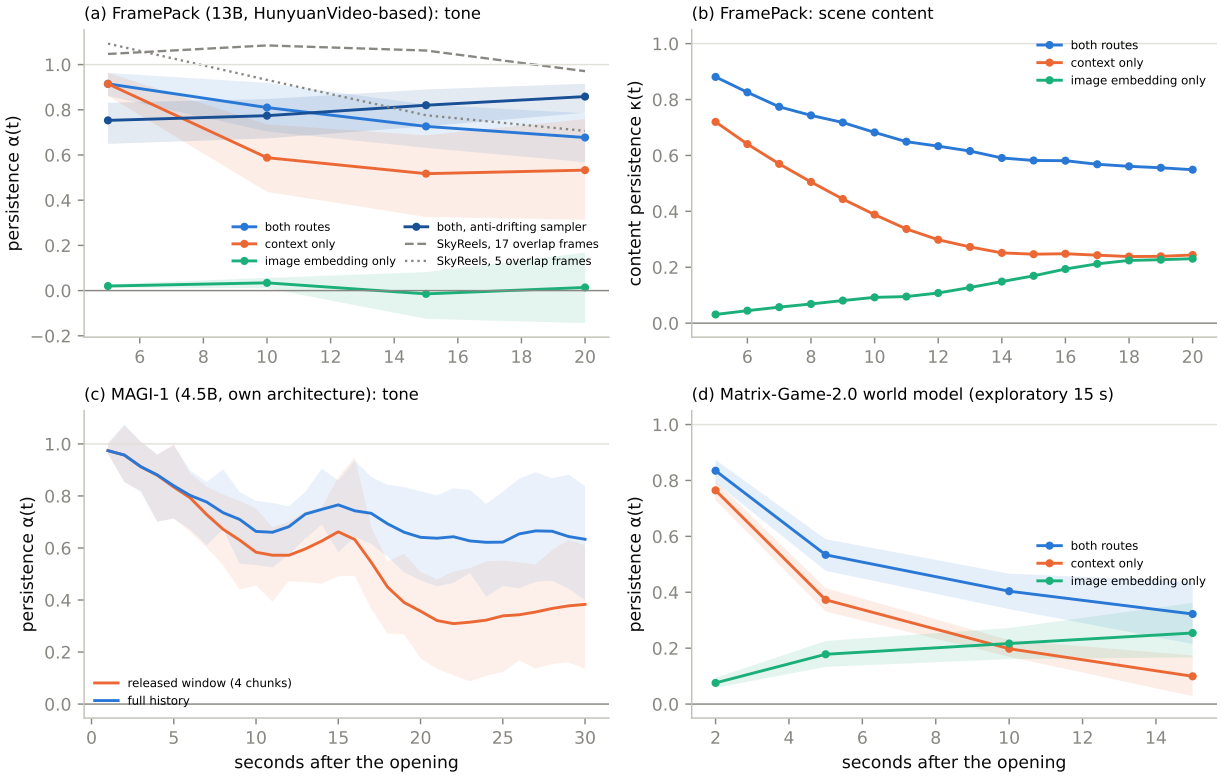


Figure 12: (a) FramePack (study 11, 12 scenes): the opening's brightness shift sent through both routes, only the clean context prefix, or only the SigLIP image embedding; the anti-drifting sampler; SkyReels-V2 with 17 and 5 overlapping frames for reference. (b) FramePack, scene content (study 17, 24 openings, city prompt): with a city frame in the context, the opening's embedding alone pulls the video toward the opening's class. (c) MAGI-1 (study 18, 12 scenes): the released window of four 1-s chunks against full history (1-s windows, 3-s moving average). (d) Matrix-Game-2.0, an action-conditioned world model without text (study 15, exploratory 15-s analysis, 20 scenes): the context latent leaves the window and is forgotten, while the rollout drifts toward the CLIP-embedded opening. Bands: scene-bootstrap 95% CIs.

distance within the range seen in training) does not reduce D (-0.03 [$-0.79, 0.72$]) and increases K ($+0.36$). Darkening persists with a sink (last-2-s $z_L = -1.99$). D and K rise roughly linearly over 31 s under every strategy, without a visible plateau.

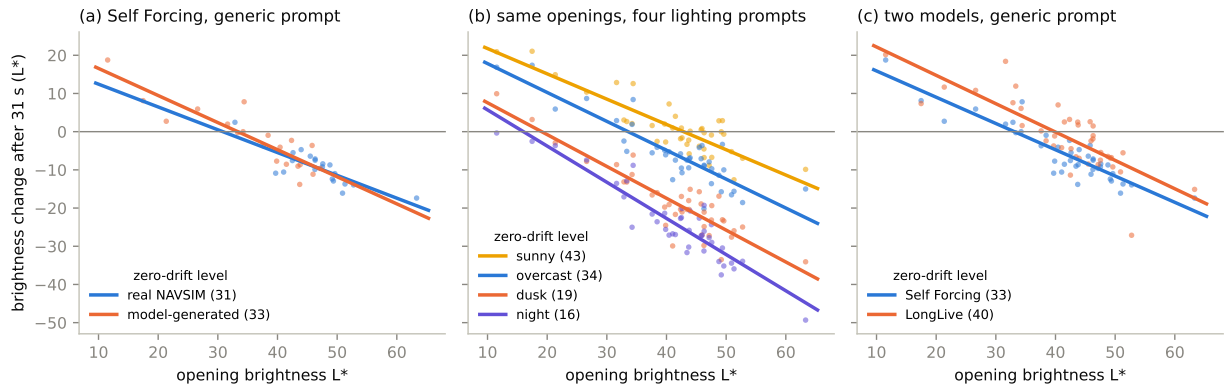


Figure 13: Brightness change after 31 s against the opening's brightness (one dot per opening, generic or lighting prompt): (a) real and generated openings follow one line; (b) four lighting prompts set four defaults (Self Forcing); (c) two models: similar slopes, different defaults.